

Group16_Project_Final_Report

May 3, 2025

```
[ ]: #For Final Project with an objective to find a good model that is able to predict which water pumps are functional and which are not.
```

```
[31]: #All relevant packages are imported
```

```
import pandas as pd
import numpy as np
import seaborn as sns
import missingno as msno
from scipy import stats
import matplotlib.pyplot as plt
from sklearn.metrics import log_loss
from imblearn.over_sampling import SMOTE
from sklearn.feature_selection import RFE
from tensorflow.keras.regularizers import l2
from tensorflow.keras.models import Sequential
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from tensorflow.keras.layers import Dense, Dropout, Input
from sklearn.model_selection import cross_val_score, GridSearchCV
from sklearn.preprocessing import StandardScaler, PolynomialFeatures
from sklearn.naive_bayes import GaussianNB, MultinomialNB, BernoulliNB
from sklearn.metrics import confusion_matrix, auc, roc_auc_score, roc_curve, classification_report, log_loss, accuracy_score, ConfusionMatrixDisplay, RocCurveDisplay, precision_score, recall_score, f1_score
from sklearn.feature_selection import SelectFromModel
from sklearn.impute import KNNImputer
from sklearn.ensemble import BaggingClassifier
import tensorflow as tf
from tensorflow import keras
from tensorflow.keras import layers, callbacks, regularizers
from sklearn.metrics import log_loss
import os

# Set seed for reproducibility
```

```

SEED = 42
import os, random, numpy as np
random.seed(SEED)                      # Python's built-in random module
np.random.seed(SEED)                    # NumPy's random module
os.environ['PYTHONHASHSEED'] = str(SEED) # Hash seed (affects dict ordering)
                                         #etc. in rare cases)

df = pd.read_csv(r"C:\Users\HP\Desktop\2025SPRING\DSCI5240\PROJECT STATUS\REPORT\DSCI 5240 Project Data.csv")
print(df.head())

```

```

import warnings
warnings.filterwarnings("ignore", category=FutureWarning)

```

| | Water Pump ID | Water Source | Type | Water Quality | Distance to Nearest Town |
|---|---------------|--------------|-------|---------------|--------------------------|
| 0 | WP001 | Well | Clean | | 44.0 |
| 1 | WP002 | Lake | Clean | | 13.0 |
| 2 | WP003 | Lake | Clean | | 27.0 |
| 3 | WP004 | Well | Clean | | 14.0 |
| 4 | WP005 | Lake | Clean | | 41.0 |

| | Population Served | Installation Year | Funder | Payment Type |
|---|-------------------|-------------------|------------|--------------|
| 0 | 13000.0 | 2006.0 | World Bank | Free |
| 1 | 13000.0 | 1990.0 | Red Cross | Free |
| 2 | 12000.0 | 1997.0 | Oxfam | Pay per use |
| 3 | 9000.0 | 1992.0 | Oxfam | Pay per use |
| 4 | 16000.0 | 2006.0 | NaN | Pay per use |

| | Water Pump Age | Pump Type | GPS Coordinates |
|---|----------------|----------------|---|
| 0 | 18.0 | Motorized Pump | (-20.599463060030295, 26.696000047794744) |
| 1 | 34.0 | Hand Pump | (-20.69129769992364, 23.313405231404484) |
| 2 | 27.0 | Hand Pump | (-19.830951420391948, 26.650358442338003) |
| 3 | 32.0 | NaN | (-22.335866062765565, 22.83485684389231) |
| 4 | 18.0 | Hand Pump | (-21.099305692773278, 24.799143614430015) |

| | Functioning Status |
|---|--------------------|
| 0 | Functioning |
| 1 | Not Functioning |
| 2 | Not Functioning |
| 3 | Functioning |
| 4 | Functioning |

[32]: #Data visualization and exploratory analysis codes have been removed from this python file so as to make it easier for model codes to run.
#Exploratory Data Analysis has already been submitted and shared in Python code file in Project Status Report.

```
#This file contains only the relevant data cleaning codes precisely required
for running all algorithms sufficiently well.
```

```
# DATA CLEANING STARTS
```

[33]: # MISSING DATA IMPUTATION: Numerical Features

```
# Taking imputation action based on distribution
# Features with normal distribution - Fill missing values with mean and replace
negative values with mean for specified columns

# Fill missing values with mean and replace negative values with mean for
'Distance to Nearest Town'
mean_distance = df['Distance to Nearest Town'].mean()
df['Distance to Nearest Town'] = df['Distance to Nearest Town'].apply(lambda x:
mean_distance if x < 0 or pd.isna(x) else x)

# Fill missing values with mean for 'Population Served'
df['Population Served'] = df['Population Served'].fillna(df['Population
Served'].mean())

# Taking imputation action based on distribution
# Features without normal distribution
# Fill missing values with mode for specified columns
mode_value = df['Installation Year'].mode()[0] # Get the mode (most frequent
value)
df['Installation Year'] = df['Installation Year'].fillna(mode_value)
mode_value = df['Water Pump Age'].mode()[0] # Get the mode (most frequent
value)
df['Water Pump Age'] = df['Water Pump Age'].fillna(mode_value)

# Checking missing values again after application of imputation technique
missing_values_df = pd.DataFrame({'Feature': df.columns, 'Missing Values': df.
isnull().sum().values})
print(missing_values_df)
```

| | Feature | Missing Values |
|---|--------------------------|----------------|
| 0 | Water Pump ID | 250 |
| 1 | Water Source Type | 250 |
| 2 | Water Quality | 250 |
| 3 | Distance to Nearest Town | 0 |
| 4 | Population Served | 0 |
| 5 | Installation Year | 0 |
| 6 | Funder | 250 |
| 7 | Payment Type | 250 |
| 8 | Water Pump Age | 0 |
| 9 | Pump Type | 250 |

| | | |
|----|--------------------|-----|
| 10 | GPS Coordinates | 250 |
| 11 | Functioning Status | 250 |

[34]: # Missing Values Check

```
# Checking the count of number of rows with missing values in any column, after
# applying imputation techniques on numerical/float data type features
missing_rows = df.isnull().any(axis=1).sum()
print(f"Number of rows with at least one missing value: {missing_rows}")
```

Number of rows with at least one missing value: 1667

[35]: # MISSING DATA IMPUTATION: Categorical Features

```
# Proportional imputation function
def proportional_imputation(df, categorical_cols):
    """
    Imputes missing values in categorical columns of a DataFrame,
    preserving the original proportions of each category within the column.

    Args:
        df (pd.DataFrame): The DataFrame to impute.
        categorical_cols (list): A list of column names that are categorical.

    Returns:
        pd.DataFrame: The DataFrame with missing values imputed.
    """
    for col in categorical_cols:
        # Calculate the existing value counts and their proportions
        value_counts = df[col].value_counts(normalize=True)

        # Identify the missing values in the column
        missing_mask = df[col].isnull()
        num_missing = missing_mask.sum()

        # If there are no missing values, skip to the next column
        if num_missing == 0:
            continue

        # Randomly choose values to fill the missing spots based on the
        # proportions
        imputed_values = np.random.choice(value_counts.index, size=num_missing,
                                           p=value_counts.values)

        # Fill the missing values with the randomly chosen values
        df.loc[missing_mask, col] = imputed_values

    return df # Return the modified DataFrame
```

```

# List of categorical columns to apply proportional imputation
categorical_features = ['Water Source Type', 'Water Quality',
                       'Funder', 'Payment Type', 'Pump Type',
                       'Functioning Status']

# Apply proportional imputation directly to the original DataFrame
df = proportional_imputation(df, categorical_features)

# Check for missing values after imputation
print("Missing values after imputation:")
print(df.isnull().sum()) # Prints the number of missing values for each column

```

Missing values after imputation:

| | |
|--------------------------|-----|
| Water Pump ID | 250 |
| Water Source Type | 0 |
| Water Quality | 0 |
| Distance to Nearest Town | 0 |
| Population Served | 0 |
| Installation Year | 0 |
| Funder | 0 |
| Payment Type | 0 |
| Water Pump Age | 0 |
| Pump Type | 0 |
| GPS Coordinates | 250 |
| Functioning Status | 0 |

dtype: int64

[36]: # Missing Values Check again

```

missing_values_df = pd.DataFrame({'Feature': df.columns, 'Missing Values': df.
                                   ↪isnull().sum().values})
print(missing_values_df)

```

| | Feature | Missing Values |
|----|--------------------------|----------------|
| 0 | Water Pump ID | 250 |
| 1 | Water Source Type | 0 |
| 2 | Water Quality | 0 |
| 3 | Distance to Nearest Town | 0 |
| 4 | Population Served | 0 |
| 5 | Installation Year | 0 |
| 6 | Funder | 0 |
| 7 | Payment Type | 0 |
| 8 | Water Pump Age | 0 |
| 9 | Pump Type | 0 |
| 10 | GPS Coordinates | 250 |
| 11 | Functioning Status | 0 |

[37]: #Handling Outliers

```
# Calculate IQR for the 'Population Served' column
Q1 = df['Population Served'].quantile(0.25)
Q3 = df['Population Served'].quantile(0.75)
IQR = Q3 - Q1

# Calculate lower and upper bounds for outliers
lower_bound = Q1 - 1.5 * IQR
upper_bound = Q3 + 1.5 * IQR

# Identify outliers (values below lower_bound or above upper_bound)
outliers = df[(df['Population Served'] < lower_bound) | (df['Population Served'] > upper_bound)]

# Option 1: Remove outliers
df_no_outliers = df[~df['Population Served'].isin(outliers['Population Served'])]

# Option 2: Cap outliers (replace outliers with lower/upper bounds)
df['Population Served'] = df['Population Served'].clip(lower=lower_bound, upper=upper_bound)

# Print the result
print("Outliers removed:")
print(outliers)
print("\nData after handling outliers:")
print(df.head())
```

Outliers removed:

| | Water Pump ID | Water Source | Type | Water Quality | Distance to Nearest Town | \ |
|------|---------------|--------------|--------------|---------------|--------------------------|---|
| 351 | WP352 | | Lake | Clean | 20.0 | |
| 1623 | WP1624 | | Well | Clean | 37.0 | |
| 1761 | WP1762 | | Lake | Clean | 50.0 | |
| 1782 | WP1783 | | Well | Clean | 48.0 | |
| 2114 | WP2115 | | Lake | Clean | 6.0 | |
| 2391 | WP2392 | | Well | Clean | 34.0 | |
| 2503 | WP2504 | | Well | Clean | 35.0 | |
| 2706 | WP2707 | | Lake | Clean | 44.0 | |
| 2932 | WP2933 | | Lake | Clean | 45.0 | |
| 3072 | WP3073 | | Well | Contaminated | 13.0 | |
| 3207 | WP3208 | River | Contaminated | | 46.0 | |
| 3518 | WP3519 | | Lake | Contaminated | 12.0 | |
| 3534 | WP3535 | | Lake | Clean | 36.0 | |
| 3684 | WP3685 | | Lake | Contaminated | 15.0 | |
| 4277 | WP4278 | | Well | Contaminated | 7.0 | |
| 4609 | WP4610 | | Lake | Clean | 40.0 | |

| | | | | |
|------|--------|------|--------------|------|
| 4691 | WP4692 | Lake | Clean | 35.0 |
| 4766 | WP4767 | Lake | Clean | 28.0 |
| 4913 | WP4914 | Well | Contaminated | 37.0 |

| | Population Served | Installation Year | Funder | Payment Type | \ |
|------|-------------------|-------------------|-----------|--------------|---|
| 351 | 22000.0 | 2020.0 | Red Cross | Free | |
| 1623 | 4000.0 | 2006.0 | Oxfam | Pay per use | |
| 1761 | 22000.0 | 2014.0 | Red Cross | Pay per use | |
| 1782 | 22000.0 | 1998.0 | UNICEF | Pay per use | |
| 2114 | 22000.0 | 1991.0 | USAID | Pay per use | |
| 2391 | 4000.0 | 2012.0 | Red Cross | Pay per use | |
| 2503 | 3000.0 | 2001.0 | Oxfam | Pay per use | |
| 2706 | 22000.0 | 2004.0 | UNICEF | Pay per use | |
| 2932 | 4000.0 | 2006.0 | Red Cross | Pay per use | |
| 3072 | 4000.0 | 2006.0 | UNICEF | Free | |
| 3207 | 2000.0 | 1998.0 | Red Cross | Pay per use | |
| 3518 | 4000.0 | 2011.0 | Red Cross | Pay per use | |
| 3534 | 4000.0 | 2006.0 | Oxfam | Pay per use | |
| 3684 | 4000.0 | 2014.0 | Red Cross | Free | |
| 4277 | 22000.0 | 2012.0 | USAID | Pay per use | |
| 4609 | 4000.0 | 2016.0 | USAID | Pay per use | |
| 4691 | 22000.0 | 1993.0 | Red Cross | Pay per use | |
| 4766 | 22000.0 | 1990.0 | Oxfam | Free | |
| 4913 | 4000.0 | 2014.0 | Red Cross | Pay per use | |

| | Water Pump Age | Pump Type | \ |
|------|----------------|----------------|---|
| 351 | 13.0 | Hand Pump | |
| 1623 | 18.0 | Motorized Pump | |
| 1761 | 10.0 | Motorized Pump | |
| 1782 | 26.0 | Motorized Pump | |
| 2114 | 33.0 | Motorized Pump | |
| 2391 | 12.0 | Hand Pump | |
| 2503 | 23.0 | Hand Pump | |
| 2706 | 20.0 | Hand Pump | |
| 2932 | 18.0 | Motorized Pump | |
| 3072 | 18.0 | Hand Pump | |
| 3207 | 26.0 | Hand Pump | |
| 3518 | 28.0 | Motorized Pump | |
| 3534 | 18.0 | Hand Pump | |
| 3684 | 13.0 | Hand Pump | |
| 4277 | 12.0 | Solar Pump | |
| 4609 | 13.0 | Motorized Pump | |
| 4691 | 31.0 | Hand Pump | |
| 4766 | 34.0 | Motorized Pump | |
| 4913 | 10.0 | Motorized Pump | |

| | | |
|-----|--|--------------------|
| | GPS Coordinates | Functioning Status |
| 351 | (-18.93340276075529, 25.208132833432785) | Not Functioning |

| | | |
|------|---|-----------------|
| 1623 | (-22.5011154948989, 22.44232179713561) | Functioning |
| 1761 | (-19.91705627831145, 22.88334295790117) | Not Functioning |
| 1782 | (-21.395185744722475, 21.83400215915583) | Functioning |
| 2114 | (-18.445136023192276, 19.83568408473501) | Not Functioning |
| 2391 | NaN | Not Functioning |
| 2503 | (-21.846849904094423, 27.875202606048454) | Functioning |
| 2706 | (-20.115143959613338, 27.96842372307367) | Not Functioning |
| 2932 | (-23.880332866455312, 22.718900901753344) | Functioning |
| 3072 | (-21.947956908132944, 23.705769697466252) | Not Functioning |
| 3207 | (-23.82235931998412, 22.517421974546835) | Functioning |
| 3518 | (-22.3381539369204, 22.574595841673172) | Not Functioning |
| 3534 | (-22.93903359350564, 22.475883440760576) | Not Functioning |
| 3684 | (-22.014944464080397, 23.970354016215058) | Not Functioning |
| 4277 | (-19.175938524367503, 28.088288834906066) | Functioning |
| 4609 | (-18.829727335050492, 24.234630718789543) | Not Functioning |
| 4691 | (-20.885178466819852, 26.412262104843453) | Not Functioning |
| 4766 | (-21.689982610996893, 26.041732729255877) | Not Functioning |
| 4913 | (-21.307601730750495, 25.1211502927296) | Not Functioning |

Data after handling outliers:

| | Water Pump ID | Water Source | Type | Water Quality | Distance to Nearest Town | \ |
|---|---------------|--------------|------|---------------|--------------------------|---|
| 0 | WP001 | | Well | Clean | 44.0 | |
| 1 | WP002 | | Lake | Clean | 13.0 | |
| 2 | WP003 | | Lake | Clean | 27.0 | |
| 3 | WP004 | | Well | Clean | 14.0 | |
| 4 | WP005 | | Lake | Clean | 41.0 | |

| | Population Served | Installation Year | Funder | Payment Type | \ |
|---|-------------------|-------------------|------------|--------------|---|
| 0 | 13000.0 | 2006.0 | World Bank | Free | |
| 1 | 13000.0 | 1990.0 | Red Cross | Free | |
| 2 | 12000.0 | 1997.0 | Oxfam | Pay per use | |
| 3 | 9000.0 | 1992.0 | Oxfam | Pay per use | |
| 4 | 16000.0 | 2006.0 | USAID | Pay per use | |

| | Water Pump Age | Pump Type | GPS Coordinates | \ |
|---|----------------|----------------|---|---|
| 0 | 18.0 | Motorized Pump | (-20.599463060030295, 26.696000047794744) | |
| 1 | 34.0 | Hand Pump | (-20.69129769992364, 23.313405231404484) | |
| 2 | 27.0 | Hand Pump | (-19.830951420391948, 26.650358442338003) | |
| 3 | 32.0 | Hand Pump | (-22.335866062765565, 22.83485684389231) | |
| 4 | 18.0 | Hand Pump | (-21.099305692773278, 24.799143614430015) | |

| | Functioning Status |
|---|--------------------|
| 0 | Functioning |
| 1 | Not Functioning |
| 2 | Not Functioning |
| 3 | Functioning |
| 4 | Functioning |

```
[38]: # FIXING INCONSISTENCIES

#Convert Installation Year to Integer
df['Installation Year'] = df['Installation Year'].astype(int)

#Fix GPS Coordinates (Split into Latitude & Longitude)
df[['Latitude', 'Longitude']] = df['GPS Coordinates'].str.extract(r'\\((.*), (.+*))\\').astype(float)

#Checking the dataset top rows again after correcting inconsistencies:
print(df.head())
```

| | Water Pump ID | Water Source | Type | Water Quality | Distance to Nearest Town |
|---|---------------|--------------|------|---------------|--------------------------|
| 0 | WP001 | | Well | Clean | 44.0 |
| 1 | WP002 | | Lake | Clean | 13.0 |
| 2 | WP003 | | Lake | Clean | 27.0 |
| 3 | WP004 | | Well | Clean | 14.0 |
| 4 | WP005 | | Lake | Clean | 41.0 |

| | Population Served | Installation Year | Funder | Payment Type |
|---|-------------------|-------------------|------------|--------------|
| 0 | 13000.0 | 2006 | World Bank | Free |
| 1 | 13000.0 | 1990 | Red Cross | Free |
| 2 | 12000.0 | 1997 | Oxfam | Pay per use |
| 3 | 9000.0 | 1992 | Oxfam | Pay per use |
| 4 | 16000.0 | 2006 | USAID | Pay per use |

| | Water Pump Age | Pump Type | GPS Coordinates |
|---|----------------|----------------|---|
| 0 | 18.0 | Motorized Pump | (-20.599463060030295, 26.696000047794744) |
| 1 | 34.0 | Hand Pump | (-20.69129769992364, 23.313405231404484) |
| 2 | 27.0 | Hand Pump | (-19.830951420391948, 26.650358442338003) |
| 3 | 32.0 | Hand Pump | (-22.335866062765565, 22.83485684389231) |
| 4 | 18.0 | Hand Pump | (-21.099305692773278, 24.799143614430015) |

| | Functioning Status | Latitude | Longitude |
|---|--------------------|------------|-----------|
| 0 | Functioning | -20.599463 | 26.696000 |
| 1 | Not Functioning | -20.691298 | 23.313405 |
| 2 | Not Functioning | -19.830951 | 26.650358 |
| 3 | Functioning | -22.335866 | 22.834857 |
| 4 | Functioning | -21.099306 | 24.799144 |

```
[39]: # Drop unnecessary columns
df = df.drop(['Water Pump ID', 'GPS Coordinates'], axis=1)
```

```
[40]: # MISSING DATA IMPUTATION:Remaining Features
```

```
#Imputation using KNNImputer of the remaining features
```

```

# Select relevant numeric columns (those that might relate to location)
numeric_cols = ['Latitude', 'Longitude', 'Distance to Nearest Town',
                 'Population Served', 'Installation Year', 'Water Pump Age']

# Create a subset of the dataframe
df_numeric = df[numeric_cols]

# Initialize KNNImputer (k=5 is common)
imputer = KNNImputer(n_neighbors=5)

# Fit and transform the data
df_imputed = imputer.fit_transform(df_numeric)

# Create back a DataFrame and assign imputed Lat/Long to original df
df_imputed = pd.DataFrame(df_imputed, columns=numeric_cols)

# Replace original Latitude and Longitude with imputed ones
df['Latitude'] = df_imputed['Latitude']
df['Longitude'] = df_imputed['Longitude']

```

[41]: #ONE HOT ENCODING

```

# One-hot encode categorical variables (excluding target for now)
df_encoded = pd.get_dummies(df, columns=['Water Source Type', 'Water Quality',
                                           'Funder', 'Payment Type', 'Pump_Type'],
                             drop_first=True)

# Encode target column (Functioning: 1, Not Functioning: 0)
df_encoded['Functioning Status'] = df['Functioning Status'].map({'Functioning': 1,
                                                               'Not Functioning': 0})

```

[42]: #Due to performance of Models not improving, we decided to drop irrelevant features, which were earlier not dropped

```

# Since water pump age and Installation Year are strongly negatively correlated, it is reasonable to keep only one to reduce collinearity. One feature is enough to tell about the other.

# List of columns to drop
columns_to_drop_nn = [
    'Latitude', 'Longitude',
    'Funder_Red Cross', 'Funder_UNICEF', 'Funder_USAID', 'Funder_World Bank',
    'Payment Type_Pay per use',
    'Installation Year'
]

# Drop the specified columns from df_encoded

```

```

df_encoded = df_encoded.drop(columns=columns_to_drop_nn)

# Check to confirm they're gone
print("\nUpdated Column Names:\n", df_encoded.columns.tolist())

```

Updated Column Names:

- ['Distance to Nearest Town', 'Population Served', 'Water Pump Age', 'Functioning Status', 'Water Source Type_Lake', 'Water Source Type_River', 'Water Source Type_Well', 'Water Quality_Contaminated', 'Pump Type_Motorized Pump', 'Pump Type_Solar Pump']

[43]: # PREPARAING FOR MODELLING

```

# Define feature columns and target
X = df_encoded.drop(columns=['Functioning Status'])
y = df_encoded['Functioning Status']

# Split first
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3,
                                                    random_state=42, stratify=y)

# Apply SMOTE to training set
smote = SMOTE(random_state=42)
X_res, y_res = smote.fit_resample(X_train, y_train)

# Scale numeric columns
scaler = StandardScaler()
numeric_cols = ['Distance to Nearest Town', 'Population Served', 'Water PumpAge']
X_res[numeric_cols] = scaler.fit_transform(X_res[numeric_cols])
X_test[numeric_cols] = scaler.transform(X_test[numeric_cols])

```

[44]: # MODELING STARTS;

```

# For every Algorithm, first a baseline algorithm was ran and later versions
# with improved parameters
# Logistic Regression (6) , Decision Tree (5), Ensemble Method (3), Naive Bayes
# (3), Neural Network (3)

```

[45]: # Logistic Regression 1: Basic Logistic Regression

```

# Set seed for reproducibility
SEED = 42
random.seed(SEED)
np.random.seed(SEED)
os.environ['PYTHONHASHSEED'] = str(SEED)

```

```

# Basic Logistic Regression with random_state
model1 = LogisticRegression(max_iter=1000, random_state=SEED)
model1.fit(X_res, y_res)

train_acc1 = model1.score(X_res, y_res)
test_acc1 = model1.score(X_test, y_test)
y_test_pred1 = model1.predict(X_test)
y_pred_proba1 = model1.predict_proba(X_test)

print("Basic Logistic Regression")
print(f"Train Accuracy: {train_acc1:.4f}")
print(f"Test Accuracy: {test_acc1:.4f}")

# Create a figure with 1 row and 2 columns
fig, axes = plt.subplots(1, 2, figsize=(14, 6))

# Plot Confusion Matrix on the first subplot
cm = confusion_matrix(y_test, y_test_pred1)
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=["0", "1"], yticklabels=["0", "1"],
            ax=axes[0])
axes[0].set_title("Confusion Matrix")
axes[0].set_xlabel("Predicted")
axes[0].set_ylabel("Actual")

# Plot ROC Curve on the second subplot
roc_auc = roc_auc_score(y_test, y_pred_proba1[:, 1])
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba1[:, 1])
axes[1].plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (area =\n{roc_auc:.2f})')
axes[1].plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
axes[1].set_xlabel('False Positive Rate')
axes[1].set_ylabel('True Positive Rate')
axes[1].set_title('ROC Curve')
axes[1].legend(loc='lower right')

# Adjust layout
plt.tight_layout()
plt.show()

print("\nClassification Report (Test Data):\n")
print(classification_report(y_test, y_test_pred1))


logloss1 = log_loss(y_test, y_pred_proba1)
print(f"Log Loss: {logloss1:.4f}")

```

Basic Logistic Regression

Train Accuracy: 0.7453

Test Accuracy: 0.7213



Classification Report (Test Data):

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.77 | 0.74 | 0.76 | 881 |
| 1 | 0.65 | 0.69 | 0.67 | 619 |
| accuracy | | | 0.72 | 1500 |
| macro avg | 0.71 | 0.72 | 0.71 | 1500 |
| weighted avg | 0.72 | 0.72 | 0.72 | 1500 |

Log Loss: 0.5688

[46]:


Logistic Regression with Hyperparameter Tuning (Model 6)

Best Params: {'C': 1, 'penalty': 'l2', 'solver': 'liblinear'}

Train Accuracy: 0.7460

Test Accuracy: 0.7200

Log Loss: 0.5687



Classification Report (Test Data):

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.77 | 0.74 | 0.76 | 881 |
| 1 | 0.65 | 0.69 | 0.67 | 619 |
| accuracy | | | 0.72 | 1500 |
| macro avg | 0.71 | 0.72 | 0.71 | 1500 |
| weighted avg | 0.72 | 0.72 | 0.72 | 1500 |


[47]:

```
Logistic Regression with Bagging
Train Accuracy: 0.7496
Test Accuracy: 0.7287
```

Classification Report (Test Data):

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.71 | 0.78 | 0.74 | 618 |
| 1 | 0.75 | 0.68 | 0.71 | 617 |
| accuracy | | | 0.73 | 1235 |
| macro avg | 0.73 | 0.73 | 0.73 | 1235 |
| weighted avg | 0.73 | 0.73 | 0.73 | 1235 |

AUC: 0.79



[48]: #Logistic Regression 2: with Regularization (with random state)

```
model2 = LogisticRegression(penalty='l2', C=1.0, solver='lbfgs', max_iter=1000, random_state=42)
model2.fit(X_res, y_res)

train_acc2 = model2.score(X_res, y_res)
test_acc2 = model2.score(X_test, y_test)

cv_acc2 = cross_val_score(model2, X_res, y_res, cv=5, scoring='accuracy').mean()

# Predictions
y_test_pred2 = model2.predict(X_test)
y_pred_proba2 = model2.predict_proba(X_test)

# Log Loss
logloss2 = log_loss(y_test, y_pred_proba2)

print("\nLogistic Regression with Regularization (Model 2)")
print(f"Train Accuracy: {train_acc2:.4f}")
print(f"Test Accuracy: {test_acc2:.4f}")
print(f"Cross-Validated Accuracy: {cv_acc2:.4f}")
print(f"Log Loss: {logloss2:.4f}")

# Evaluation Metrics
cm = confusion_matrix(y_test, y_test_pred2)
roc_auc = roc_auc_score(y_test, y_pred_proba2[:, 1])
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba2[:, 1])

fig, axes = plt.subplots(1, 2, figsize=(12, 5))
```

```

sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=["0", "1"], yticklabels=["0", "1"], ax=axes[0])
axes[0].set_title("Confusion Matrix (Model 2)")
axes[0].set_xlabel("Predicted")
axes[0].set_ylabel("Actual")

axes[1].plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (area = {roc_auc:.2f})')
axes[1].plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
axes[1].set_xlabel('False Positive Rate')
axes[1].set_ylabel('True Positive Rate')
axes[1].set_title('ROC Curve (Model 2)')
axes[1].legend(loc='lower right')

plt.tight_layout()
plt.show()

print("\nClassification Report (Test Data):\n")
print(classification_report(y_test, y_test_pred2))

```


Logistic Regression with Regularization (Model 2)

Train Accuracy: 0.7453

Test Accuracy: 0.7344

Cross-Validated Accuracy: 0.7431

Log Loss: 0.5393



Classification Report (Test Data):

| precision | recall | f1-score | support |
|-----------|--------|----------|---------|
|-----------|--------|----------|---------|

| | | | | |
|--------------|------|------|------|------|
| 0 | 0.72 | 0.77 | 0.74 | 618 |
| 1 | 0.75 | 0.70 | 0.72 | 617 |
| accuracy | | | 0.73 | 1235 |
| macro avg | 0.74 | 0.73 | 0.73 | 1235 |
| weighted avg | 0.74 | 0.73 | 0.73 | 1235 |

[49]: #Logistic Regression 3: with RFE and SMOTE (random state added to SMOTE and model)

```
base_model = LogisticRegression(max_iter=1000, random_state=42)
selector = RFE(base_model, n_features_to_select=8)
selector.fit(X_train, y_train)

X_train_rfe = selector.transform(X_train)
X_test_rfe = selector.transform(X_test)

X_rfe_res, y_rfe_res = SMOTE(random_state=42).fit_resample(X_train_rfe, y_train)

model4 = LogisticRegression(max_iter=1000, random_state=42)
model4.fit(X_rfe_res, y_rfe_res)

train_acc4 = model4.score(X_rfe_res, y_rfe_res)
test_acc4 = model4.score(X_test_rfe, y_test)
cv_acc4 = cross_val_score(model4, X_rfe_res, y_rfe_res, cv=5, scoring='accuracy').mean()

y_test_pred4 = model4.predict(X_test_rfe)
y_pred_proba4 = model4.predict_proba(X_test_rfe)
logloss4 = log_loss(y_test, y_pred_proba4)

print("\nLogistic Regression with RFE (Model 4)")
print(f"Train Accuracy: {train_acc4:.4f}")
print(f"Test Accuracy: {test_acc4:.4f}")
print(f"Cross-Validated Accuracy: {cv_acc4:.4f}")
print(f"Log Loss: {logloss4:.4f}")

cm = confusion_matrix(y_test, y_test_pred4)
roc_auc = roc_auc_score(y_test, y_pred_proba4[:, 1])
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba4[:, 1])

fig, axes = plt.subplots(1, 2, figsize=(12, 5))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=["0", "1"], yticklabels=["0", "1"], ax=axes[0])
axes[0].set_title("Confusion Matrix (Model 4)")
axes[0].set_xlabel("Predicted")
```

```

axes[0].set_ylabel("Actual")


axes[1].plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (area = {roc_auc:.2f})')
axes[1].plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
axes[1].set_xlabel('False Positive Rate')
axes[1].set_ylabel('True Positive Rate')
axes[1].set_title('ROC Curve (Model 4)')
axes[1].legend(loc='lower right')

plt.tight_layout()
plt.show()

print("\nClassification Report (Test Data):\n")
print(classification_report(y_test, y_test_pred4))

```

Logistic Regression with RFE (Model 4)
 Train Accuracy: 0.7490
 Test Accuracy: 0.7320
 Cross-Validated Accuracy: 0.7507
 Log Loss: 0.5423



Classification Report (Test Data):

| | precision | recall | f1-score | support |
|----------|-----------|--------|----------|---------|
| 0 | 0.71 | 0.78 | 0.74 | 618 |
| 1 | 0.75 | 0.69 | 0.72 | 617 |
| accuracy | | | 0.73 | 1235 |

| | | | | |
|--------------|------|------|------|------|
| macro avg | 0.73 | 0.73 | 0.73 | 1235 |
| weighted avg | 0.73 | 0.73 | 0.73 | 1235 |

[50]: #Logistic Regression 4: with Polynomial Features (random state added to model)

```

poly = PolynomialFeatures(degree=2, include_bias=False)
X_poly_train = poly.fit_transform(X_res)
X_poly_test = poly.transform(X_test)

model5 = LogisticRegression(max_iter=1000, random_state=42)
model5.fit(X_poly_train, y_res)

train_acc5 = model5.score(X_poly_train, y_res)
test_acc5 = model5.score(X_poly_test, y_test)

y_test_pred5 = model5.predict(X_poly_test)
y_pred_proba5 = model5.predict_proba(X_poly_test)
logloss5 = log_loss(y_test, y_pred_proba5)

print("\nLogistic Regression with Polynomial Features (Model 5)")
print(f"Train Accuracy: {train_acc5:.4f}")
print(f"Test Accuracy: {test_acc5:.4f}")
print(f"Log Loss: {logloss5:.4f}")

cm = confusion_matrix(y_test, y_test_pred5)
roc_auc = roc_auc_score(y_test, y_pred_proba5[:, 1])
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba5[:, 1])

fig, axes = plt.subplots(1, 2, figsize=(12, 5))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=["0", "1"], yticklabels=["0", "1"], ax=axes[0])
axes[0].set_title("Confusion Matrix (Model 5)")
axes[0].set_xlabel("Predicted")
axes[0].set_ylabel("Actual")

axes[1].plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (area = {roc_auc:.2f})')
axes[1].plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
axes[1].set_xlabel('False Positive Rate')
axes[1].set_ylabel('True Positive Rate')
axes[1].set_title('ROC Curve (Model 5)')
axes[1].legend(loc='lower right')

plt.tight_layout()
plt.show()

```


```
print("\nClassification Report (Test Data):\n")
print(classification_report(y_test, y_test_pred5))
```

Logistic Regression with Polynomial Features (Model 5)

Train Accuracy: 0.7489

Test Accuracy: 0.7466

Log Loss: 0.5305



Classification Report (Test Data):

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.73 | 0.78 | 0.75 | 618 |
| 1 | 0.76 | 0.71 | 0.74 | 617 |
| accuracy | | | 0.75 | 1235 |
| macro avg | 0.75 | 0.75 | 0.75 | 1235 |
| weighted avg | 0.75 | 0.75 | 0.75 | 1235 |

```
[ ]: # Logistic Regression 5: with Hyperparameter Tuning (GridSearchCV)
param_grid = {
    'C': [0.01, 0.1, 1, 10],
    'penalty': ['l1', 'l2'],
    'solver': ['liblinear']
}
grid_model = GridSearchCV(LogisticRegression(max_iter=1000, random_state=42), param_grid, cv=5)
grid_model.fit(X_res, y_res)
```

```

train_acc6 = grid_model.score(X_res, y_res)
test_acc6 = grid_model.score(X_test, y_test)

y_test_pred6 = grid_model.predict(X_test)
y_pred_proba6 = grid_model.predict_proba(X_test)
logloss6 = log_loss(y_test, y_pred_proba6)

print("\nLogistic Regression with Hyperparameter Tuning (Model 6)")
print(f"Best Params: {grid_model.best_params_}")
print(f"Train Accuracy: {train_acc6:.4f}")
print(f"Test Accuracy: {test_acc6:.4f}")
print(f"Log Loss: {logloss6:.4f}")

cm = confusion_matrix(y_test, y_test_pred6)
roc_auc = roc_auc_score(y_test, y_pred_proba6[:, 1])
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba6[:, 1])

fig, axes = plt.subplots(1, 2, figsize=(12, 5))
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", xticklabels=["0", "1"], yticklabels=["0", "1"], ax=axes[0])
axes[0].set_title("Confusion Matrix (Model 6)")
axes[0].set_xlabel("Predicted")
axes[0].set_ylabel("Actual")

axes[1].plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (area = {roc_auc:.2f})')
axes[1].plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
axes[1].set_xlabel('False Positive Rate')
axes[1].set_ylabel('True Positive Rate')
axes[1].set_title('ROC Curve (Model 6)')
axes[1].legend(loc='lower right')

plt.tight_layout()
plt.show()

print("\nClassification Report (Test Data):\n")
print(classification_report(y_test, y_test_pred6))

```

[51]: #Logistic Regression 6: with Common Important Features

```

#Feature Importance Intersection
# Set random seed for reproducibility
np.random.seed(42)

# Feature importance
base_model = LogisticRegression(max_iter=1000, solver='liblinear', random_state=42)
rfe = RFE(base_model, n_features_to_select=8)

```

```

rfe.fit(X_train, y_train)
rfe_features = X_train.columns[rfe.support_]

l1_model = LogisticRegression(penalty='l1', solver='liblinear', C=1,□
    ↪max_iter=1000, random_state=42)
l1_model.fit(X_train, y_train)
l1_features = X_train.columns[(l1_model.coef_[0] != 0)]

basic_model = LogisticRegression(max_iter=1000, random_state=42)
basic_model.fit(X_train, y_train)
coeff_basic = pd.Series(basic_model.coef_[0], index=X_train.columns)
top_basic = coeff_basic.abs().sort_values(ascending=False).head(10).index

common_features = set(rfe_features) & set(l1_features) & set(top_basic)
print("Common important features across RFE, L1, and Coefficients:\n",□
    ↪common_features)

# Plot
palette = sns.color_palette("crest", n_colors=6)
sorted_coeffs = coeff_basic.abs().sort_values(ascending=False).head(10)

feature_importance_df = pd.DataFrame({
    'Feature': sorted_coeffs.index,
    'Importance': sorted_coeffs.values,
    'Common': sorted_coeffs.index.isin(common_features)
})
custom_colors = {True: palette[4], False: palette[1]}
plt.figure(figsize=(10, 6))
sns.barplot(
    data=feature_importance_df,
    x='Importance',
    y='Feature',
    hue='Common',
    dodge=False,
    palette=custom_colors
)
plt.title('Top 10 Feature Importances (Logistic Regression Coefficients)',□
    ↪fontsize=14)
plt.xlabel('Absolute Coefficient Value')
plt.ylabel('Feature')
plt.legend(title='Common Across RFE, L1, Coeffs')
plt.tight_layout()
plt.show()

# Select only the common important features
X_train_selected = X_train[list(common_features)]
X_test_selected = X_test[list(common_features)]

```

```

# Build the model
model_selected = LogisticRegression(max_iter=1000, solver='liblinear', random_state=42)
model_selected.fit(X_train_selected, y_train)

# Predict on train and test data
y_train_pred = model_selected.predict(X_train_selected)
y_test_pred = model_selected.predict(X_test_selected)
y_pred_proba_selected = model_selected.predict_proba(X_test_selected)

# Evaluate
train_accuracy = accuracy_score(y_train, y_train_pred)
test_accuracy = accuracy_score(y_test, y_test_pred)
logloss_selected = log_loss(y_test, y_pred_proba_selected)

print("\nLogistic Regression with Selected Features")
print(f"Train Accuracy (Selected Features): {train_accuracy:.4f}")
print(f"Test Accuracy (Selected Features): {test_accuracy:.4f}")
print(f"Log Loss: {logloss_selected:.4f}")

# Metrics
cm = confusion_matrix(y_test, y_test_pred)
roc_auc = roc_auc_score(y_test, y_pred_proba_selected[:, 1])
fpr, tpr, thresholds = roc_curve(y_test, y_pred_proba_selected[:, 1])

# Set up side-by-side plots
fig, axes = plt.subplots(1, 2, figsize=(12, 5))

# Confusion Matrix
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues",
            xticklabels=["0", "1"],
            yticklabels=["0", "1"],
            ax=axes[0])
axes[0].set_title("Confusion Matrix")
axes[0].set_xlabel("Predicted")
axes[0].set_ylabel("Actual")

# ROC Curve
axes[1].plot(fpr, tpr, color='darkorange', lw=2, label=f'ROC curve (area = {roc_auc:.2f})')
axes[1].plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
axes[1].set_xlabel('False Positive Rate')
axes[1].set_ylabel('True Positive Rate')
axes[1].set_title('ROC Curve')
axes[1].legend(loc='lower right')

```

```

plt.tight_layout()
plt.show()

print("\nClassification Report (Test Data):\n")
print(classification_report(y_test, y_test_pred))


```

Common important features across RFE, L1, and Coefficients:

```

{'Population Served', 'Water Source Type_River', 'Distance to Nearest Town',
'Pump Type_Motorized Pump', 'Pump Type_Solar Pump', 'Water Source Type_Well',
'Water Source Type_Lake', 'Water Quality_Contaminated'}

```




Logistic Regression with Selected Features

Train Accuracy (Selected Features): 0.7492

Test Accuracy (Selected Features): 0.7312

Log Loss: 0.5426



Classification Report (Test Data) :

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.71 | 0.77 | 0.74 | 618 |
| 1 | 0.75 | 0.69 | 0.72 | 617 |
| accuracy | | | 0.73 | 1235 |
| macro avg | 0.73 | 0.73 | 0.73 | 1235 |
| weighted avg | 0.73 | 0.73 | 0.73 | 1235 |

[52]: #Decision Tree Modelling

[53]: # Define a function to evaluate models

```
def evaluate_model(model, X_train, y_train, X_test, y_test, model_name="Model"):
    print(f"\n== {model_name} ==")

    # Predictions
    y_pred_train = model.predict(X_train)
    y_pred_test = model.predict(X_test)

    # Probabilities for ROC
    if hasattr(model, "predict_proba"):
        y_probs_test = model.predict_proba(X_test)[:, 1]
    else:
        y_probs_test = None

    # Accuracy
```

```

print(f"Train Accuracy: {accuracy_score(y_train, y_pred_train):.4f}")
print(f"Test Accuracy: {accuracy_score(y_test, y_pred_test):.4f}")

# Precision, Recall, F1-Score
print("\nClassification Report (Test Data):")
print(classification_report(y_test, y_pred_test))

# ROC-AUC
if y_probs_test is not None:
    roc_auc = roc_auc_score(y_test, y_probs_test)
    print(f"ROC-AUC Score (Test Data): {roc_auc:.4f}")

# 5-fold Cross Validation Accuracy
cv_scores = cross_val_score(model, X_train, y_train, cv=5,
                             scoring='accuracy')
print(f"Cross-Validation Accuracy (mean ± std): {cv_scores.mean():.4f} ± {cv_scores.std():.4f}")

# Plot Confusion Matrix and ROC Curve side-by-side
fig, axes = plt.subplots(1, 2, figsize=(12, 5))

# Confusion Matrix
cm = confusion_matrix(y_test, y_pred_test)
disp = ConfusionMatrixDisplay(confusion_matrix=cm)
disp.plot(ax=axes[0], values_format='d')
axes[0].set_title(f"{model_name} - Confusion Matrix")

# ROC Curve
if y_probs_test is not None:
    fpr, tpr, _ = roc_curve(y_test, y_probs_test)
    RocCurveDisplay(fpr=fpr, tpr=tpr).plot(ax=axes[1])
    axes[1].set_title(f"{model_name} - ROC Curve")

plt.tight_layout()
plt.show()

```

[54]: # Model 1 Basic Decision Tree

```

print("\n==== DECISION TREE v1 - BASIC ====")
dt1 = DecisionTreeClassifier(random_state=42)
dt1.fit(X_train, y_train)
evaluate_model(dt1, X_train, y_train, X_test, y_test, model_name="Decision Tree v1 (Basic)")

# Model 2 Pruned Decision Tree (Max_Depth=5):
print("\n==== DECISION TREE Max_Depth=5 ====")
dt = DecisionTreeClassifier(max_depth=5, random_state=42)
dt.fit(X_train, y_train)

```

```

evaluate_model(dt, X_train, y_train, X_test, y_test, model_name="Decision Tree
↳(max_depth=5)")

# Model 3 Optimized Tree Depth (Max_Depth=6)
print("\n==== DECISION TREE v2 - Tuned max_depth=6 ===")
dt2 = DecisionTreeClassifier(max_depth=6, random_state=42)
dt2.fit(X_train, y_train)
evaluate_model(dt2, X_train, y_train, X_test, y_test, model_name="Decision Tree
↳v2 (max_depth=6)")

# Model 4 Criterion and Minimum Samples Split Optimization (Entropy, depth=8,
↳min_samples_split=10)
print("\n==== DECISION TREE v3 - Entropy, Depth=8, Min Split=10 ===")
dt3 = DecisionTreeClassifier(criterion='entropy', max_depth=8,
↳min_samples_split=10, random_state=42)
dt3.fit(X_train, y_train)
evaluate_model(dt3, X_train, y_train, X_test, y_test, model_name="Decision Tree
↳v3 (Entropy, depth=8, split=10)")

# Model 5 Tuned Decision Tree (SMOTE + Feature Selection)
print("\n==== DECISION TREE - Tuned with SMOTE + Feature Selection ===")
# SMOTE resampling
smote = SMOTE(random_state=42)
X_train_res, y_train_res = smote.fit_resample(X_train, y_train)

# Feature selection
selector = SelectFromModel(RandomForestClassifier(n_estimators=100,
↳random_state=42))
selector.fit(X_train_res, y_train_res)
X_train_sel = selector.transform(X_train_res)
X_test_sel = selector.transform(X_test)

# Hyperparameter tuning
dt_params = {
    'criterion': ['gini', 'entropy'],
    'max_depth': [6, 8, 10, None],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 3, 5],
    'class_weight': [None, 'balanced']
}
dt_grid = GridSearchCV(DecisionTreeClassifier(random_state=42), dt_params,
↳cv=5, scoring='accuracy')
dt_grid.fit(X_train_sel, y_train_res)

dt_best = dt_grid.best_estimator_
print("\nBest Parameters from Grid Search:", dt_grid.best_params_)

```

```
evaluate_model(dt_best, X_train_sel, y_train_res, X_test_sel, y_test, model_name="Tuned Decision Tree (GridSearch)")
```

==== DECISION TREE Max_Depth=5 ===

==== Decision Tree (max_depth=5) ===

Train Accuracy: 0.7471

Test Accuracy: 0.7174

Classification Report (Test Data):


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.68 | 0.84 | 0.75 | 618 |
| 1 | 0.79 | 0.60 | 0.68 | 617 |
| accuracy | | | 0.72 | 1235 |
| macro avg | 0.73 | 0.72 | 0.71 | 1235 |
| weighted avg | 0.73 | 0.72 | 0.71 | 1235 |

ROC-AUC Score (Test Data): 0.7828

Cross-Validation Accuracy (mean ± std): 0.7339 ± 0.0198

C:\Users\HP\anaconda3\Lib\site-packages\sklearn\metrics_plot\roc_curve.py:189:
UserWarning: No artists with labels found to put in legend. Note that artists
whose label start with an underscore are ignored when legend() is called with no
argument.

```
self.ax_.legend(loc="lower right")
```



==== DECISION TREE v1 - BASIC ===

```
==== Decision Tree v1 (Basic) ====
```

```
Train Accuracy: 0.9951
```

```
Test Accuracy: 0.6688
```

```
Classification Report (Test Data):
```


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.67 | 0.66 | 0.67 | 618 |
| 1 | 0.66 | 0.68 | 0.67 | 617 |
| accuracy | | | 0.67 | 1235 |
| macro avg | 0.67 | 0.67 | 0.67 | 1235 |
| weighted avg | 0.67 | 0.67 | 0.67 | 1235 |

```
ROC-AUC Score (Test Data): 0.6686
```

```
Cross-Validation Accuracy (mean ± std): 0.6488 ± 0.0188
```

```
C:\Users\HP\anaconda3\Lib\site-packages\sklearn\metrics\_plot\roc_curve.py:189:  
UserWarning: No artists with labels found to put in legend. Note that artists  
whose label start with an underscore are ignored when legend() is called with no  
argument.
```

```
    self.ax_.legend(loc="lower right")
```



```
==== DECISION TREE v2 - Tuned max_depth=6 ====
```

```
==== Decision Tree v2 (max_depth=6) ====
```

```
Train Accuracy: 0.7614
```

```
Test Accuracy: 0.7352
```

Classification Report (Test Data):


| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.72 | 0.76 | 0.74 | 618 |
| 1 | 0.75 | 0.71 | 0.73 | 617 |
| accuracy | | | 0.74 | 1235 |
| macro avg | 0.74 | 0.74 | 0.74 | 1235 |
| weighted avg | 0.74 | 0.74 | 0.74 | 1235 |

ROC-AUC Score (Test Data): 0.7848

Cross-Validation Accuracy (mean ± std): 0.7343 ± 0.0238

C:\Users\HP\anaconda3\Lib\site-packages\sklearn\metrics_plot\roc_curve.py:189:
UserWarning: No artists with labels found to put in legend. Note that artists
whose label start with an underscore are ignored when legend() is called with no
argument.

```
self.ax_.legend(loc="lower right")
```



==== DECISION TREE v3 - Entropy, Depth=8, Min Split=10 ===

==== Decision Tree v3 (Entropy, depth=8, split=10) ===

Train Accuracy: 0.7714

Test Accuracy: 0.7328

Classification Report (Test Data):

| | precision | recall | f1-score | support |
|---|-----------|--------|----------|---------|
| 0 | 0.73 | 0.74 | 0.74 | 618 |
| 1 | 0.74 | 0.72 | 0.73 | 617 |


| | | | | |
|--------------|------|------|------|------|
| accuracy | | | 0.73 | 1235 |
| macro avg | 0.73 | 0.73 | 0.73 | 1235 |
| weighted avg | 0.73 | 0.73 | 0.73 | 1235 |

ROC-AUC Score (Test Data): 0.7784

Cross-Validation Accuracy (mean ± std): 0.7367 ± 0.0191

C:\Users\HP\anaconda3\Lib\site-packages\sklearn\metrics_plot\roc_curve.py:189:
UserWarning: No artists with labels found to put in legend. Note that artists
whose label start with an underscore are ignored when legend() is called with no
argument.

```
    self.ax_.legend(loc="lower right")
```



Best Parameters from Grid Search: {'class_weight': None, 'criterion': 'entropy', 'max_depth': 6, 'min_samples_leaf': 3, 'min_samples_split': 10}

==== Tuned Decision Tree (GridSearch) ===

Train Accuracy: 0.7347

Test Accuracy: 0.7036

Classification Report (Test Data):


| | precision | recall | f1-score | support |
|-----------|-----------|--------|----------|---------|
| 0 | 0.68 | 0.78 | 0.73 | 618 |
| 1 | 0.74 | 0.62 | 0.68 | 617 |
| accuracy | | | 0.70 | 1235 |
| macro avg | 0.71 | 0.70 | 0.70 | 1235 |

```

weighted avg      0.71      0.70      0.70      1235

ROC-AUC Score (Test Data): 0.7370
Cross-Validation Accuracy (mean ± std): 0.7267 ± 0.0186
C:\Users\HP\anaconda3\Lib\site-packages\sklearn\metrics\_plot\roc_curve.py:189:
UserWarning: No artists with labels found to put in legend. Note that artists
whose label start with an underscore are ignored when legend() is called with no
argument.
    self.ax_.legend(loc="lower right")

```



```
[55]: #Ensemble Method
```

```

[56]: print("\n==== RANDOM FOREST MODELS ====")

# Helper function to evaluate models
def evaluate_model(model, X_train, X_test, y_train, y_test):
    train_acc = model.score(X_train, y_train)
    test_acc = model.score(X_test, y_test)
    overfit_gap = train_acc - test_acc
    y_pred = model.predict(X_test)
    y_proba = model.predict_proba(X_test)[:,1]

# Cross-validation
cv_scores = cross_val_score(model, X_train, y_train, cv=5)

# Metrics
roc_auc = roc_auc_score(y_test, y_proba)
f1 = f1_score(y_test, y_pred)
cm = confusion_matrix(y_test, y_pred)
cr = classification_report(y_test, y_pred)

```

```
# Print Metrics
print(f"Train Accuracy: {train_acc:.4f}")
print(f"Test Accuracy: {test_acc:.4f}")
print(f"Overfitting Gap (Train - Test): {overfit_gap:.4f}")
print(f"Cross-Validation Score: {cv_scores.mean():.4f} ± {cv_scores.std():.4f}")
print(f"ROC-AUC Score: {roc_auc:.4f}")
print(f"F1-Score: {f1:.4f}")
print("\nClassification Report:\n", cr)

# Side-by-side Plots
fig, ax = plt.subplots(1, 2, figsize=(14, 5))

# Confusion Matrix
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", ax=ax[0])
ax[0].set_title('Confusion Matrix')
ax[0].set_xlabel('Predicted')
ax[0].set_ylabel('Actual')

# ROC Curve
fpr, tpr, thresholds = roc_curve(y_test, y_proba)
ax[1].plot(fpr, tpr, color='darkorange')
ax[1].plot([0, 1], [0, 1], color='navy', linestyle='--')
ax[1].set_title('ROC Curve')
ax[1].set_xlabel('False Positive Rate')
ax[1].set_ylabel('True Positive Rate')

plt.tight_layout()
plt.show()
```

==== RANDOM FOREST MODELS ====

[57] : # ENSEMBLE METHOD

```
# Random Forest 1: Basic Random Forest
rf1 = RandomForestClassifier(random_state=42)
rf1.fit(X_train, y_train)
print("\nRandom Forest v1 - Basic:")
evaluate_model(rf1, X_train, X_test, y_train, y_test)

# Random Forest 2: With Constrained Random Forest (Tuned with n_estimators=100, max_depth=10)
rf2 = RandomForestClassifier(n_estimators=100, max_depth=10, random_state=42)
rf2.fit(X_train, y_train)
print("\nRandom Forest v2 - n_estimators=100, max_depth=10:")
```

```
evaluate_model(rf2, X_train, X_test, y_train, y_test)
```

Random Forest v1 - Basic:

Train Accuracy: 0.9951

Test Accuracy: 0.7166

Overfitting Gap (Train - Test): 0.2785


Cross-Validation Score: 0.7221 ± 0.0139

ROC-AUC Score: 0.7830

F1-Score: 0.7131

Classification Report:

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.71 | 0.73 | 0.72 | 618 |
| 1 | 0.72 | 0.71 | 0.71 | 617 |
| accuracy | | | 0.72 | 1235 |
| macro avg | 0.72 | 0.72 | 0.72 | 1235 |
| weighted avg | 0.72 | 0.72 | 0.72 | 1235 |



Random Forest v2 - n_estimators=100, max_depth=10:

Train Accuracy: 0.8229

Test Accuracy: 0.7352

Overfitting Gap (Train - Test): 0.0876

Cross-Validation Score: 0.7454 ± 0.0183


ROC-AUC Score: 0.7934

F1-Score: 0.7282

Classification Report:

| | precision | recall | f1-score | support |
|--|-----------|--------|----------|---------|
|--|-----------|--------|----------|---------|

| | | | | |
|--------------|------|------|------|------|
| 0 | 0.72 | 0.76 | 0.74 | 618 |
| 1 | 0.75 | 0.71 | 0.73 | 617 |
| accuracy | | | 0.74 | 1235 |
| macro avg | 0.74 | 0.74 | 0.74 | 1235 |
| weighted avg | 0.74 | 0.74 | 0.74 | 1235 |



```
[ ]: # Ensemble Method 3

#Bagging with Logistic Regression(Random Seed)

# Set random seed to ensure reproducibility
random_seed = 42

# Ensure consistent train-test split with random_state
X_train, X_test, y_train, y_test = train_test_split(X_res, y_res, test_size=0.
    ↪3, random_state=random_seed)

# Initialize the base model
base_model = LogisticRegression(random_state=random_seed)

# Create the bagging ensemble with the random_state parameter
bagging_model = BaggingClassifier(estimator=base_model, n_estimators=10, ↪
    ↪random_state=random_seed)

# Fit on the resampled training data
bagging_model.fit(X_train, y_train)

# Predictions
y_train_pred = bagging_model.predict(X_train)
y_test_pred = bagging_model.predict(X_test)
```

```

# Calculate probabilities for ROC curve
y_test_proba = bagging_model.predict_proba(X_test)[:, 1]

# Performance metrics
train_accuracy = accuracy_score(y_train, y_train_pred)
test_accuracy = accuracy_score(y_test, y_test_pred)

print(f"Logistic Regression with Bagging")
print(f"Train Accuracy: {train_accuracy:.4f}")
print(f"Test Accuracy: {test_accuracy:.4f}")

# Confusion matrix
cm = confusion_matrix(y_test, y_test_pred)

# Classification report
print("\nClassification Report (Test Data):")
print(classification_report(y_test, y_test_pred))

# ROC curve calculation
fpr, tpr, _ = roc_curve(y_test, y_test_proba)
roc_auc = auc(fpr, tpr)
print(f"AUC: {roc_auc:.2f}")

# Plot Confusion Matrix and ROC Curve side-by-side
fig, axes = plt.subplots(1, 2, figsize=(12, 5))

# Confusion Matrix
sns.heatmap(cm, annot=True, fmt="d", cmap="Blues", ax=axes[0])
axes[0].set_title('Confusion Matrix')
axes[0].set_xlabel('Predicted')
axes[0].set_ylabel('Actual')

# ROC Curve
axes[1].plot(fpr, tpr, label=f'AUC = {roc_auc:.2f}')
axes[1].plot([0, 1], [0, 1], 'k--') # Random guess line
axes[1].set_xlim([0.0, 1.0])
axes[1].set_ylim([0.0, 1.05])
axes[1].set_xlabel('False Positive Rate')
axes[1].set_ylabel('True Positive Rate')
axes[1].set_title('ROC Curve')
axes[1].legend(loc="lower right")

plt.tight_layout()
plt.show()

```

[58]: #Naive Bayes

```
[59]: # Model 1: Gaussian Naive Bayes

# Gaussian Naive Bayes assumes features are continuous and normally distributed
# (Gaussian).

# Initialize the GaussianNB model and Train the model using the training data
gnb = GaussianNB()
gnb.fit(X_res, y_res) # Corrected to use X_res and y_res

# Predict outcomes on training and testing datasets
y_train_pred = gnb.predict(X_res)
y_test_pred = gnb.predict(X_test)

# Evaluate performance by checking training and testing accuracies
train_acc_gnb = accuracy_score(y_res, y_train_pred)
test_acc_gnb = accuracy_score(y_test, y_test_pred)

print(f"Training Accuracy (GaussianNB): {train_acc_gnb:.4f}")
print(f"Testing Accuracy (GaussianNB): {test_acc_gnb:.4f}")

# Overfitting check
if abs(train_acc_gnb - test_acc_gnb) < 0.05:
    print("No major overfitting detected for GaussianNB.")
else:
    print("Possible overfitting detected for GaussianNB.")

# To visualize performance:
# Confusion Matrix and ROC Curve Side-by-Side
fig, axes = plt.subplots(ncols=2, figsize=(12, 5))

# Confusion Matrix
cm = confusion_matrix(y_test, y_test_pred)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=np.unique(y_test_pred),
            yticklabels=np.unique(y_test), ax=axes[0])
axes[0].set_title("GaussianNB Confusion Matrix")
axes[0].set_xlabel("Predicted")
axes[0].set_ylabel("Actual")

# ROC Curve
y_test_binary = y_test.astype(int)

y_probs = gnb.predict_proba(X_test)[:, 1]
fpr, tpr, _ = roc_curve(y_test_binary, y_probs)

axes[1].plot(fpr, tpr, color='blue', label='GaussianNB')
axes[1].plot([0, 1], [0, 1], 'k--')
```

```


axes[1].set_title("ROC Curve - GaussianNB")
axes[1].set_xlabel("False Positive Rate")
axes[1].set_ylabel("True Positive Rate")
axes[1].legend()

plt.tight_layout()
plt.show()

print("AUC Score (GaussianNB):", roc_auc_score(y_test_binary, y_probs))

```

Training Accuracy (GaussianNB): 0.7253
 Testing Accuracy (GaussianNB): 0.7166
 No major overfitting detected for GaussianNB.



AUC Score (GaussianNB): 0.7476055976040241

```

[60]: # Method 2. GaussianNB Improvement: Manual Top Feature Selection
# -----
# Step 1: Map 'Functioning' and 'Not Functioning' to binary (0/1)
y_binary = y.map({'Not Functioning': 0, 'Functioning': 1})

# Temporarily add the binary target to feature dataframe
df_temp = X.copy()
df_temp['Functioning Status'] = y_binary.values

# Now compute correlation
correlations = df_temp.corr()['Functioning Status'].abs().
    sort_values(ascending=False)

# Select top 20 features most correlated with 'Functioning Status'
top_features = correlations.index[1:21] # Exclude the target itself

```

```

# Ensure 'Functioning Status' is excluded from top_features
top_features = [feature for feature in top_features if feature != 'Functioning\u2022
    ↵Status']

print("Top 20 features selected based on highest correlation:\n", top_features)

# Step 2: Subset X_res and X_test using only the selected features
X_res_top = X_res[top_features]
X_test_top = X_test[top_features]

# Step 3: Train GaussianNB on top selected features
gnb_top = GaussianNB()
gnb_top.fit(X_res_top, y_res)

# Step 4: Predict outcomes on training and testing datasets
y_train_pred_top = gnb_top.predict(X_res_top)
y_test_pred_top = gnb_top.predict(X_test_top)

# Step 5: Evaluate performance by checking training and testing accuracies
train_acc_gnb_top = accuracy_score(y_res, y_train_pred_top)
test_acc_gnb_top = accuracy_score(y_test, y_test_pred_top)

print(f"Training Accuracy (GaussianNB with Top Features): {train_acc_gnb_top:.4f}")
print(f"Testing Accuracy (GaussianNB with Top Features): {test_acc_gnb_top:.4f}")

# Confusion Matrix and ROC Curve Side-by-Side
fig, axes = plt.subplots(ncols=2, figsize=(12, 5))

# Confusion Matrix
cm_top = confusion_matrix(y_test, y_test_pred_top)
sns.heatmap(cm_top, annot=True, fmt='d', cmap='Blues',
            xticklabels=np.unique(y_test_pred_top),
            yticklabels=np.unique(y_test), ax=axes[0])
axes[0].set_title("GNB Top Features Confusion Matrix")
axes[0].set_xlabel("Predicted")
axes[0].set_ylabel("Actual")

# ROC Curve
y_test_binary = y_test.astype(int)

y_probs_top = gnb_top.predict_proba(X_test_top)[:, 1]
fpr_top, tpr_top, _ = roc_curve(y_test_binary, y_probs_top)

axes[1].plot(fpr_top, tpr_top, color='green', label='GNB Top Features')

```

```

        axes[1].plot([0, 1], [0, 1], 'k--')
        axes[1].set_title("ROC Curve - GNB Top Features")
        axes[1].set_xlabel("False Positive Rate")
        axes[1].set_ylabel("True Positive Rate")
        axes[1].legend()

        plt.tight_layout()
        plt.show()

print("AUC Score (GNB Top Features):", roc_auc_score(y_test_binary, y_probs_top))

```

Top 20 features selected based on highest correlation:


```

['Population Served', 'Water Pump Age', 'Water Source Type_Lake', 'Water Source
Type_River', 'Water Source Type_Well', 'Water Quality_Contaminated', 'Pump
Type_Motorized Pump', 'Pump Type_Solar Pump']

```

Training Accuracy (GaussianNB with Top Features): 0.7224

Testing Accuracy (GaussianNB with Top Features): 0.7142



AUC Score (GNB Top Features): 0.7375441246662785

```

[61]: # Model 3: Multinomial Naive Bayes
#-----

# Clip negatives (MultinomialNB requires non-negative features)
X_res_mnb = X_res.clip(lower=0)
X_test_mnb = X_test.clip(lower=0)

# Train MultinomialNB
mnb = MultinomialNB()
mnb.fit(X_res_mnb, y_res)

```

```

# Predict outcomes on training and testing datasets
y_train_pred_mnb = mnb.predict(X_res_mnb)
y_test_pred_mnb = mnb.predict(X_test_mnb)

# Evaluate performance by checking training and testing accuracies
train_acc_mnb = accuracy_score(y_res, y_train_pred_mnb)
test_acc_mnb = accuracy_score(y_test, y_test_pred_mnb)

print(f"Training Accuracy (MultinomialNB): {train_acc_mnb:.4f}")
print(f"Testing Accuracy (MultinomialNB): {test_acc_mnb:.4f}")

# Overfitting check
if abs(train_acc_mnb - test_acc_mnb) < 0.05:
    print("No major overfitting detected for MultinomialNB.")
else:
    print("Possible overfitting detected for MultinomialNB.")

# Visualization: Confusion Matrix + ROC Curve Side-by-Side
#-----
fig, axes = plt.subplots(ncols=2, figsize=(12, 5))

# Confusion matrix
cm = confusion_matrix(y_test, y_test_pred_mnb)
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues',
            xticklabels=['Not Functioning', 'Functioning'],
            yticklabels=['Not Functioning', 'Functioning'],
            ax=axes[0])
axes[0].set_title("MultinomialNB Confusion Matrix")
axes[0].set_xlabel("Predicted")
axes[0].set_ylabel("Actual")

# ROC Curve
# Important: For ROC curve, make sure y_test is binary (0/1)
y_test_binary = y_test.astype(int) # Corrected

y_probs_mnb = mnb.predict_proba(X_test_mnb)[:, 1]
fpr, tpr, _ = roc_curve(y_test_binary, y_probs_mnb)


axes[1].plot(fpr, tpr, color='blue', label='MultinomialNB')
axes[1].plot([0, 1], [0, 1], 'k--')
axes[1].set_title("ROC Curve - MultinomialNB")
axes[1].set_xlabel("False Positive Rate")
axes[1].set_ylabel("True Positive Rate")
axes[1].legend()

plt.tight_layout()
plt.show()

```

```
print("AUC Score (MultinomialNB):", roc_auc_score(y_test_binary, y_probs_mnb))
```

Training Accuracy (MultinomialNB): 0.7200
Testing Accuracy (MultinomialNB): 0.7166
No major overfitting detected for MultinomialNB.



AUC Score (MultinomialNB): 0.7549238669205312

[62]: #Neural network

[63]: # 1. Regularization-Focused Neural Network

```
# 1. Define early stopping
early_stop = callbacks.EarlyStopping(
    monitor='val_loss',
    patience=10,
    restore_best_weights=True
)

# 2. Define the model
model_b = keras.Sequential([
    layers.Input(shape=(X_res.shape[1],)),
    layers.Dense(128, activation='relu', kernel_regularizer=regularizers.l2(0.01)),
    layers.BatchNormalization(),
    layers.Dropout(0.3),

    layers.Dense(64, activation='relu', kernel_regularizer=regularizers.l2(0.01)),
    layers.BatchNormalization(),
```

```

        layers.Dense(1, activation='sigmoid')
    ])

# 3. Compile the model
model_b.compile(
    optimizer=keras.optimizers.Adam(learning_rate=0.001),
    loss='binary_crossentropy',
    metrics=['accuracy', keras.metrics.AUC(name='auc')])
)

# 4. Train the model
history_b = model_b.fit(
    X_res, y_res,
    validation_data=(X_test, y_test),
    epochs=100,
    batch_size=32,
    callbacks=[early_stop],
    verbose=1
)

# 5. Evaluate the model on Train set
train_loss_b, train_accuracy_b, train_auc_b = model_b.evaluate(X_res, y_res, ▾
    verbose=0)

# 6. Evaluate the model on Test set
test_loss_b, test_accuracy_b, test_auc_b = model_b.evaluate(X_test, y_test, ▾
    verbose=0)

# 7. Print Results
print("\n--- Final Evaluation: Model -Regularization-Focused ---")
print(f"Train Loss: {train_loss_b:.4f}")
print(f"Train Accuracy: {train_accuracy_b:.4f}")
print(f"Train AUC: {train_auc_b:.4f}")

print(f"\nTest Loss: {test_loss_b:.4f}")
print(f"Test Accuracy: {test_accuracy_b:.4f}")
print(f"Test AUC: {test_auc_b:.4f}")

```

Epoch 1/100
129/129 3s 5ms/step -
accuracy: 0.6310 - auc: 0.6779 - loss: 1.5899 - val_accuracy: 0.5822 - val_auc:
0.7757 - val_loss: 1.3093
Epoch 2/100
129/129 0s 3ms/step -
accuracy: 0.7198 - auc: 0.7761 - loss: 1.1740 - val_accuracy: 0.6146 - val_auc:
0.7939 - val_loss: 1.0979
Epoch 3/100

```
129/129          0s 3ms/step -
accuracy: 0.7253 - auc: 0.7777 - loss: 1.0004 - val_accuracy: 0.6850 - val_auc:
0.7974 - val_loss: 0.9391
Epoch 4/100
129/129          0s 3ms/step -
accuracy: 0.7323 - auc: 0.7915 - loss: 0.8782 - val_accuracy: 0.7255 - val_auc:
0.8066 - val_loss: 0.8106
Epoch 5/100
129/129          0s 2ms/step -
accuracy: 0.7271 - auc: 0.7919 - loss: 0.7973 - val_accuracy: 0.7457 - val_auc:
0.8060 - val_loss: 0.7385
Epoch 6/100
129/129          0s 2ms/step -
accuracy: 0.7331 - auc: 0.7943 - loss: 0.7411 - val_accuracy: 0.7433 - val_auc:
0.8142 - val_loss: 0.6931
Epoch 7/100
129/129          0s 3ms/step -
accuracy: 0.7358 - auc: 0.7973 - loss: 0.6941 - val_accuracy: 0.7385 - val_auc:
0.8112 - val_loss: 0.6629
Epoch 8/100
129/129          0s 3ms/step -
accuracy: 0.7500 - auc: 0.8063 - loss: 0.6537 - val_accuracy: 0.7449 - val_auc:
0.8127 - val_loss: 0.6320
Epoch 9/100
129/129          0s 3ms/step -
accuracy: 0.7299 - auc: 0.7945 - loss: 0.6457 - val_accuracy: 0.7425 - val_auc:
0.8038 - val_loss: 0.6206
Epoch 10/100
129/129          0s 3ms/step -
accuracy: 0.7394 - auc: 0.7962 - loss: 0.6277 - val_accuracy: 0.7433 - val_auc:
0.8138 - val_loss: 0.5990
Epoch 11/100
129/129          0s 3ms/step -
accuracy: 0.7412 - auc: 0.8024 - loss: 0.6048 - val_accuracy: 0.7490 - val_auc:
0.8127 - val_loss: 0.5856
Epoch 12/100
129/129          0s 2ms/step -
accuracy: 0.7391 - auc: 0.8066 - loss: 0.5899 - val_accuracy: 0.7296 - val_auc:
0.7994 - val_loss: 0.5887
Epoch 13/100
129/129          0s 2ms/step -
accuracy: 0.7263 - auc: 0.7908 - loss: 0.5962 - val_accuracy: 0.7336 - val_auc:
0.7906 - val_loss: 0.5929
Epoch 14/100
129/129          0s 3ms/step -
accuracy: 0.7412 - auc: 0.8009 - loss: 0.5875 - val_accuracy: 0.7393 - val_auc:
0.8024 - val_loss: 0.5699
Epoch 15/100
```

```
129/129          0s 2ms/step -
accuracy: 0.7311 - auc: 0.7889 - loss: 0.5922 - val_accuracy: 0.7474 - val_auc:
0.8105 - val_loss: 0.5574
Epoch 16/100
129/129          0s 3ms/step -
accuracy: 0.7324 - auc: 0.7920 - loss: 0.5833 - val_accuracy: 0.7457 - val_auc:
0.8096 - val_loss: 0.5575
Epoch 17/100
129/129          0s 3ms/step -
accuracy: 0.7354 - auc: 0.7870 - loss: 0.5846 - val_accuracy: 0.7514 - val_auc:
0.8046 - val_loss: 0.5567
Epoch 18/100
129/129          0s 3ms/step -
accuracy: 0.7314 - auc: 0.7932 - loss: 0.5763 - val_accuracy: 0.7457 - val_auc:
0.8010 - val_loss: 0.5580
Epoch 19/100
129/129          0s 3ms/step -
accuracy: 0.7467 - auc: 0.8107 - loss: 0.5514 - val_accuracy: 0.7328 - val_auc:
0.8040 - val_loss: 0.5550
Epoch 20/100
129/129          0s 3ms/step -
accuracy: 0.7416 - auc: 0.8067 - loss: 0.5530 - val_accuracy: 0.7360 - val_auc:
0.7963 - val_loss: 0.5615
Epoch 21/100
129/129          0s 3ms/step -
accuracy: 0.7468 - auc: 0.8029 - loss: 0.5574 - val_accuracy: 0.7263 - val_auc:
0.7793 - val_loss: 0.5798
Epoch 22/100
129/129          0s 3ms/step -
accuracy: 0.7410 - auc: 0.8054 - loss: 0.5551 - val_accuracy: 0.7393 - val_auc:
0.7982 - val_loss: 0.5566
Epoch 23/100
129/129          0s 2ms/step -
accuracy: 0.7392 - auc: 0.8028 - loss: 0.5560 - val_accuracy: 0.7385 - val_auc:
0.8072 - val_loss: 0.5470
Epoch 24/100
129/129          0s 3ms/step -
accuracy: 0.7350 - auc: 0.8017 - loss: 0.5566 - val_accuracy: 0.7417 - val_auc:
0.7999 - val_loss: 0.5522
Epoch 25/100
129/129          0s 4ms/step -
accuracy: 0.7513 - auc: 0.8119 - loss: 0.5427 - val_accuracy: 0.7474 - val_auc:
0.8079 - val_loss: 0.5430
Epoch 26/100
129/129          0s 3ms/step -
accuracy: 0.7439 - auc: 0.8059 - loss: 0.5474 - val_accuracy: 0.7425 - val_auc:
0.8051 - val_loss: 0.5440
Epoch 27/100
```

```
129/129          0s 3ms/step -
accuracy: 0.7550 - auc: 0.8141 - loss: 0.5363 - val_accuracy: 0.7506 - val_auc:
0.8041 - val_loss: 0.5430
Epoch 28/100
129/129          0s 3ms/step -
accuracy: 0.7357 - auc: 0.8010 - loss: 0.5532 - val_accuracy: 0.7360 - val_auc:
0.7968 - val_loss: 0.5537
Epoch 29/100
129/129          0s 3ms/step -
accuracy: 0.7536 - auc: 0.8147 - loss: 0.5407 - val_accuracy: 0.7441 - val_auc:
0.8034 - val_loss: 0.5434
Epoch 30/100
129/129          0s 3ms/step -
accuracy: 0.7498 - auc: 0.8092 - loss: 0.5418 - val_accuracy: 0.7441 - val_auc:
0.8060 - val_loss: 0.5408
Epoch 31/100
129/129          0s 3ms/step -
accuracy: 0.7455 - auc: 0.8116 - loss: 0.5430 - val_accuracy: 0.7425 - val_auc:
0.8048 - val_loss: 0.5413
Epoch 32/100
129/129          0s 3ms/step -
accuracy: 0.7374 - auc: 0.7958 - loss: 0.5563 - val_accuracy: 0.7482 - val_auc:
0.8013 - val_loss: 0.5428
Epoch 33/100
129/129          0s 3ms/step -
accuracy: 0.7464 - auc: 0.7984 - loss: 0.5483 - val_accuracy: 0.7433 - val_auc:
0.7968 - val_loss: 0.5486
Epoch 34/100
129/129          0s 3ms/step -
accuracy: 0.7348 - auc: 0.7891 - loss: 0.5613 - val_accuracy: 0.7506 - val_auc:
0.8039 - val_loss: 0.5380
Epoch 35/100
129/129          0s 2ms/step -
accuracy: 0.7491 - auc: 0.8026 - loss: 0.5469 - val_accuracy: 0.7174 - val_auc:
0.7951 - val_loss: 0.5550
Epoch 36/100
129/129          0s 3ms/step -
accuracy: 0.7572 - auc: 0.8165 - loss: 0.5329 - val_accuracy: 0.7401 - val_auc:
0.8017 - val_loss: 0.5439
Epoch 37/100
129/129          0s 2ms/step -
accuracy: 0.7472 - auc: 0.8058 - loss: 0.5418 - val_accuracy: 0.7514 - val_auc:
0.8045 - val_loss: 0.5416
Epoch 38/100
129/129          0s 3ms/step -
accuracy: 0.7336 - auc: 0.7989 - loss: 0.5501 - val_accuracy: 0.7449 - val_auc:
0.8051 - val_loss: 0.5409
Epoch 39/100
```

```
129/129          0s 3ms/step -
accuracy: 0.7346 - auc: 0.7925 - loss: 0.5576 - val_accuracy: 0.7482 - val_auc:
0.8038 - val_loss: 0.5398
Epoch 40/100
129/129          0s 3ms/step -
accuracy: 0.7526 - auc: 0.8156 - loss: 0.5306 - val_accuracy: 0.7466 - val_auc:
0.8025 - val_loss: 0.5409
Epoch 41/100
129/129          0s 3ms/step -
accuracy: 0.7492 - auc: 0.8099 - loss: 0.5359 - val_accuracy: 0.7466 - val_auc:
0.7983 - val_loss: 0.5447
Epoch 42/100
129/129          0s 3ms/step -
accuracy: 0.7334 - auc: 0.7971 - loss: 0.5507 - val_accuracy: 0.7425 - val_auc:
0.8042 - val_loss: 0.5371
Epoch 43/100
129/129          0s 3ms/step -
accuracy: 0.7399 - auc: 0.7976 - loss: 0.5518 - val_accuracy: 0.7449 - val_auc:
0.8014 - val_loss: 0.5396
Epoch 44/100
129/129          0s 3ms/step -
accuracy: 0.7376 - auc: 0.8043 - loss: 0.5445 - val_accuracy: 0.7425 - val_auc:
0.8008 - val_loss: 0.5405
Epoch 45/100
129/129          0s 3ms/step -
accuracy: 0.7260 - auc: 0.7861 - loss: 0.5624 - val_accuracy: 0.7417 - val_auc:
0.8022 - val_loss: 0.5388
Epoch 46/100
129/129          0s 2ms/step -
accuracy: 0.7290 - auc: 0.7964 - loss: 0.5472 - val_accuracy: 0.7506 - val_auc:
0.8019 - val_loss: 0.5407
Epoch 47/100
129/129          0s 3ms/step -
accuracy: 0.7361 - auc: 0.8028 - loss: 0.5439 - val_accuracy: 0.7474 - val_auc:
0.7999 - val_loss: 0.5405
Epoch 48/100
129/129          0s 3ms/step -
accuracy: 0.7420 - auc: 0.8003 - loss: 0.5460 - val_accuracy: 0.7474 - val_auc:
0.8048 - val_loss: 0.5370
Epoch 49/100
129/129          0s 2ms/step -
accuracy: 0.7322 - auc: 0.7845 - loss: 0.5585 - val_accuracy: 0.7466 - val_auc:
0.8009 - val_loss: 0.5427
Epoch 50/100
129/129          0s 2ms/step -
accuracy: 0.7288 - auc: 0.7865 - loss: 0.5591 - val_accuracy: 0.7344 - val_auc:
0.8012 - val_loss: 0.5404
Epoch 51/100
```

```
129/129          0s 2ms/step -
accuracy: 0.7360 - auc: 0.7916 - loss: 0.5543 - val_accuracy: 0.7482 - val_auc:
0.8003 - val_loss: 0.5409
Epoch 52/100
129/129          0s 3ms/step -
accuracy: 0.7451 - auc: 0.8038 - loss: 0.5419 - val_accuracy: 0.7385 - val_auc:
0.8003 - val_loss: 0.5425
Epoch 53/100
129/129          0s 4ms/step -
accuracy: 0.7350 - auc: 0.8016 - loss: 0.5434 - val_accuracy: 0.7385 - val_auc:
0.7975 - val_loss: 0.5446
Epoch 54/100
129/129          0s 3ms/step -
accuracy: 0.7485 - auc: 0.8079 - loss: 0.5367 - val_accuracy: 0.7441 - val_auc:
0.7981 - val_loss: 0.5408
Epoch 55/100
129/129          0s 3ms/step -
accuracy: 0.7440 - auc: 0.8017 - loss: 0.5424 - val_accuracy: 0.7368 - val_auc:
0.8035 - val_loss: 0.5390
Epoch 56/100
129/129          0s 3ms/step -
accuracy: 0.7529 - auc: 0.8096 - loss: 0.5352 - val_accuracy: 0.7166 - val_auc:
0.7921 - val_loss: 0.5518
Epoch 57/100
129/129          0s 3ms/step -
accuracy: 0.7334 - auc: 0.8018 - loss: 0.5433 - val_accuracy: 0.7441 - val_auc:
0.8035 - val_loss: 0.5366
Epoch 58/100
129/129          0s 3ms/step -
accuracy: 0.7357 - auc: 0.7970 - loss: 0.5480 - val_accuracy: 0.7360 - val_auc:
0.8004 - val_loss: 0.5376
Epoch 59/100
129/129          0s 3ms/step -
accuracy: 0.7442 - auc: 0.8005 - loss: 0.5429 - val_accuracy: 0.7457 - val_auc:
0.7998 - val_loss: 0.5387
Epoch 60/100
129/129          0s 3ms/step -
accuracy: 0.7396 - auc: 0.7984 - loss: 0.5455 - val_accuracy: 0.7449 - val_auc:
0.7984 - val_loss: 0.5384
Epoch 61/100
129/129          0s 3ms/step -
accuracy: 0.7380 - auc: 0.7999 - loss: 0.5442 - val_accuracy: 0.7344 - val_auc:
0.7880 - val_loss: 0.5494
Epoch 62/100
129/129          0s 3ms/step -
accuracy: 0.7309 - auc: 0.8025 - loss: 0.5444 - val_accuracy: 0.7474 - val_auc:
0.8058 - val_loss: 0.5338
Epoch 63/100
```

```
129/129          0s 3ms/step -
accuracy: 0.7465 - auc: 0.8035 - loss: 0.5382 - val_accuracy: 0.7425 - val_auc:
0.7993 - val_loss: 0.5409
Epoch 64/100
129/129          0s 3ms/step -
accuracy: 0.7412 - auc: 0.7998 - loss: 0.5472 - val_accuracy: 0.7433 - val_auc:
0.8030 - val_loss: 0.5364
Epoch 65/100
129/129          0s 3ms/step -
accuracy: 0.7364 - auc: 0.7929 - loss: 0.5527 - val_accuracy: 0.7449 - val_auc:
0.8019 - val_loss: 0.5364
Epoch 66/100
129/129          0s 3ms/step -
accuracy: 0.7429 - auc: 0.8032 - loss: 0.5389 - val_accuracy: 0.7433 - val_auc:
0.8018 - val_loss: 0.5365
Epoch 67/100
129/129          0s 3ms/step -
accuracy: 0.7431 - auc: 0.8033 - loss: 0.5402 - val_accuracy: 0.7425 - val_auc:
0.8012 - val_loss: 0.5391
Epoch 68/100
129/129          0s 3ms/step -
accuracy: 0.7359 - auc: 0.8026 - loss: 0.5396 - val_accuracy: 0.7490 - val_auc:
0.7998 - val_loss: 0.5356
Epoch 69/100
129/129          0s 3ms/step -
accuracy: 0.7440 - auc: 0.8079 - loss: 0.5344 - val_accuracy: 0.7352 - val_auc:
0.7990 - val_loss: 0.5407
Epoch 70/100
129/129          0s 3ms/step -
accuracy: 0.7405 - auc: 0.7997 - loss: 0.5439 - val_accuracy: 0.7498 - val_auc:
0.8034 - val_loss: 0.5348
Epoch 71/100
129/129          0s 3ms/step -
accuracy: 0.7508 - auc: 0.8113 - loss: 0.5312 - val_accuracy: 0.7457 - val_auc:
0.8031 - val_loss: 0.5366
Epoch 72/100
129/129          0s 3ms/step -
accuracy: 0.7397 - auc: 0.8012 - loss: 0.5421 - val_accuracy: 0.7474 - val_auc:
0.8009 - val_loss: 0.5374
```

--- Final Evaluation: Model -Regularization-Focused ---

```
Train Loss: 0.5279
Train Accuracy: 0.7482
Train AUC: 0.8147
```

```
Test Loss: 0.5338
Test Accuracy: 0.7474
Test AUC: 0.8058
```

```
[64]: # 2. High-Dropout Stability Neural Network

# 1. Define early stopping to prevent overfitting
early_stop = callbacks.EarlyStopping(
    monitor='val_loss',
    patience=15, # Increased patience for better stability
    restore_best_weights=True
)

# 2. Define the model with adjusted architecture
model = keras.Sequential([
    layers.Input(shape=(X_res.shape[1],)),

    # First hidden layer with L2 regularization and increased dropout
    layers.Dense(64, activation='relu', kernel_regularizer=regularizers.l2(0.
    ↪005)),
    layers.BatchNormalization(),
    layers.Dropout(0.5), # Increased dropout to reduce overfitting

    # Second hidden layer with L2 regularization and increased dropout
    layers.Dense(32, activation='relu', kernel_regularizer=regularizers.l2(0.
    ↪005)),
    layers.BatchNormalization(),
    layers.Dropout(0.4), # Increased dropout to reduce overfitting

    # Output layer
    layers.Dense(1, activation='sigmoid') # Single output neuron for binary
    ↪classification
])

# 3. Compile the model with a slightly reduced learning rate
model.compile(
    optimizer=keras.optimizers.Adam(learning_rate=0.0005), # Slightly reduced
    ↪learning rate
    loss='binary_crossentropy',
    metrics=['accuracy', keras.metrics.AUC(name='auc')]
)

# 4. Train the model
history = model.fit(
    X_res, y_res,
    validation_data=(X_test, y_test),
    epochs=100,
    batch_size=32,
    callbacks=[early_stop],
    verbose=1
)
```

```

# 5. Evaluate the model on Train set
train_loss, train_accuracy, train_auc = model.evaluate(X_res, y_res, verbose=0)

# 6. Evaluate the model on Test set
test_loss, test_accuracy, test_auc = model.evaluate(X_test, y_test, verbose=0)

# 7. Print Results
print("\n--- Final Evaluation -High-Dropout Stability ---")
print(f"Train Loss: {train_loss:.4f}")
print(f"Train Accuracy: {train_accuracy:.4f}")
print(f"Train AUC: {train_auc:.4f}")

print(f"\nTest Loss: {test_loss:.4f}")
print(f"Test Accuracy: {test_accuracy:.4f}")
print(f"Test AUC: {test_auc:.4f}")

```

Epoch 1/100
129/129 3s 5ms/step -
accuracy: 0.5315 - auc: 0.5482 - loss: 1.1245 - val_accuracy: 0.6939 - val_auc:
0.7310 - val_loss: 0.8994
Epoch 2/100
129/129 0s 3ms/step -
accuracy: 0.6278 - auc: 0.6757 - loss: 0.9631 - val_accuracy: 0.7101 - val_auc:
0.7555 - val_loss: 0.8441
Epoch 3/100
129/129 0s 3ms/step -
accuracy: 0.6625 - auc: 0.7127 - loss: 0.8988 - val_accuracy: 0.7166 - val_auc:
0.7649 - val_loss: 0.8074
Epoch 4/100
129/129 0s 3ms/step -
accuracy: 0.6626 - auc: 0.7093 - loss: 0.8804 - val_accuracy: 0.7296 - val_auc:
0.7718 - val_loss: 0.7825
Epoch 5/100
129/129 0s 3ms/step -
accuracy: 0.6826 - auc: 0.7382 - loss: 0.8315 - val_accuracy: 0.7271 - val_auc:
0.7751 - val_loss: 0.7643
Epoch 6/100
129/129 0s 3ms/step -
accuracy: 0.6876 - auc: 0.7400 - loss: 0.8141 - val_accuracy: 0.7304 - val_auc:
0.7767 - val_loss: 0.7485
Epoch 7/100
129/129 0s 3ms/step -
accuracy: 0.7061 - auc: 0.7460 - loss: 0.7923 - val_accuracy: 0.7344 - val_auc:
0.7807 - val_loss: 0.7339
Epoch 8/100
129/129 0s 3ms/step -
accuracy: 0.6943 - auc: 0.7544 - loss: 0.7692 - val_accuracy: 0.7344 - val_auc:

```
0.7821 - val_loss: 0.7215
Epoch 9/100
129/129          0s 3ms/step -
accuracy: 0.7002 - auc: 0.7537 - loss: 0.7630 - val_accuracy: 0.7377 - val_auc:
0.7839 - val_loss: 0.7081
Epoch 10/100
129/129          0s 3ms/step -
accuracy: 0.6967 - auc: 0.7460 - loss: 0.7597 - val_accuracy: 0.7360 - val_auc:
0.7843 - val_loss: 0.6978
Epoch 11/100
129/129          0s 3ms/step -
accuracy: 0.7062 - auc: 0.7596 - loss: 0.7291 - val_accuracy: 0.7344 - val_auc:
0.7866 - val_loss: 0.6865
Epoch 12/100
129/129          0s 3ms/step -
accuracy: 0.7189 - auc: 0.7762 - loss: 0.7071 - val_accuracy: 0.7385 - val_auc:
0.7888 - val_loss: 0.6758
Epoch 13/100
129/129          0s 2ms/step -
accuracy: 0.7301 - auc: 0.7807 - loss: 0.6916 - val_accuracy: 0.7409 - val_auc:
0.7902 - val_loss: 0.6659
Epoch 14/100
129/129          0s 2ms/step -
accuracy: 0.7278 - auc: 0.7836 - loss: 0.6794 - val_accuracy: 0.7433 - val_auc:
0.7932 - val_loss: 0.6555
Epoch 15/100
129/129          0s 2ms/step -
accuracy: 0.7216 - auc: 0.7761 - loss: 0.6789 - val_accuracy: 0.7433 - val_auc:
0.7929 - val_loss: 0.6481
Epoch 16/100
129/129          0s 2ms/step -
accuracy: 0.7332 - auc: 0.7879 - loss: 0.6595 - val_accuracy: 0.7433 - val_auc:
0.7955 - val_loss: 0.6390
Epoch 17/100
129/129          0s 2ms/step -
accuracy: 0.7261 - auc: 0.7783 - loss: 0.6616 - val_accuracy: 0.7417 - val_auc:
0.7968 - val_loss: 0.6317
Epoch 18/100
129/129          0s 2ms/step -
accuracy: 0.7239 - auc: 0.7809 - loss: 0.6547 - val_accuracy: 0.7449 - val_auc:
0.7949 - val_loss: 0.6251
Epoch 19/100
129/129          0s 2ms/step -
accuracy: 0.7355 - auc: 0.7887 - loss: 0.6393 - val_accuracy: 0.7457 - val_auc:
0.7953 - val_loss: 0.6188
Epoch 20/100
129/129          0s 2ms/step -
accuracy: 0.7262 - auc: 0.7843 - loss: 0.6384 - val_accuracy: 0.7433 - val_auc:
```

```
0.7972 - val_loss: 0.6118
Epoch 21/100
129/129          0s 2ms/step -
accuracy: 0.7322 - auc: 0.7934 - loss: 0.6199 - val_accuracy: 0.7425 - val_auc:
0.7989 - val_loss: 0.6062
Epoch 22/100
129/129          0s 3ms/step -
accuracy: 0.7277 - auc: 0.7957 - loss: 0.6134 - val_accuracy: 0.7417 - val_auc:
0.7984 - val_loss: 0.6003
Epoch 23/100
129/129          0s 2ms/step -
accuracy: 0.7435 - auc: 0.7936 - loss: 0.6100 - val_accuracy: 0.7417 - val_auc:
0.7987 - val_loss: 0.5947
Epoch 24/100
129/129          0s 3ms/step -
accuracy: 0.7425 - auc: 0.7968 - loss: 0.6025 - val_accuracy: 0.7417 - val_auc:
0.7998 - val_loss: 0.5897
Epoch 25/100
129/129          0s 3ms/step -
accuracy: 0.7310 - auc: 0.7953 - loss: 0.5988 - val_accuracy: 0.7385 - val_auc:
0.8001 - val_loss: 0.5851
Epoch 26/100
129/129          1s 4ms/step -
accuracy: 0.7250 - auc: 0.7841 - loss: 0.6096 - val_accuracy: 0.7449 - val_auc:
0.7995 - val_loss: 0.5816
Epoch 27/100
129/129          0s 3ms/step -
accuracy: 0.7332 - auc: 0.7939 - loss: 0.5980 - val_accuracy: 0.7433 - val_auc:
0.8003 - val_loss: 0.5764
Epoch 28/100
129/129          0s 3ms/step -
accuracy: 0.7446 - auc: 0.8025 - loss: 0.5808 - val_accuracy: 0.7441 - val_auc:
0.8010 - val_loss: 0.5735
Epoch 29/100
129/129          0s 3ms/step -
accuracy: 0.7318 - auc: 0.7878 - loss: 0.5963 - val_accuracy: 0.7466 - val_auc:
0.8005 - val_loss: 0.5711
Epoch 30/100
129/129          0s 3ms/step -
accuracy: 0.7345 - auc: 0.7938 - loss: 0.5838 - val_accuracy: 0.7425 - val_auc:
0.7972 - val_loss: 0.5715
Epoch 31/100
129/129          0s 3ms/step -
accuracy: 0.7241 - auc: 0.7966 - loss: 0.5778 - val_accuracy: 0.7474 - val_auc:
0.8014 - val_loss: 0.5653
Epoch 32/100
129/129          0s 3ms/step -
accuracy: 0.7381 - auc: 0.7900 - loss: 0.5830 - val_accuracy: 0.7466 - val_auc:
```

```
0.8000 - val_loss: 0.5622
Epoch 33/100
129/129      0s 3ms/step -
accuracy: 0.7343 - auc: 0.7989 - loss: 0.5696 - val_accuracy: 0.7433 - val_auc:
0.8013 - val_loss: 0.5590
Epoch 34/100
129/129      0s 3ms/step -
accuracy: 0.7459 - auc: 0.7996 - loss: 0.5681 - val_accuracy: 0.7457 - val_auc:
0.7972 - val_loss: 0.5608
Epoch 35/100
129/129      0s 3ms/step -
accuracy: 0.7332 - auc: 0.7902 - loss: 0.5809 - val_accuracy: 0.7449 - val_auc:
0.8021 - val_loss: 0.5563
Epoch 36/100
129/129      0s 2ms/step -
accuracy: 0.7370 - auc: 0.8022 - loss: 0.5620 - val_accuracy: 0.7490 - val_auc:
0.8011 - val_loss: 0.5535
Epoch 37/100
129/129      0s 2ms/step -
accuracy: 0.7423 - auc: 0.7956 - loss: 0.5704 - val_accuracy: 0.7466 - val_auc:
0.8002 - val_loss: 0.5538
Epoch 38/100
129/129      0s 3ms/step -
accuracy: 0.7351 - auc: 0.7983 - loss: 0.5651 - val_accuracy: 0.7441 - val_auc:
0.8015 - val_loss: 0.5526
Epoch 39/100
129/129      0s 3ms/step -
accuracy: 0.7363 - auc: 0.7953 - loss: 0.5670 - val_accuracy: 0.7401 - val_auc:
0.8010 - val_loss: 0.5514
Epoch 40/100
129/129      0s 3ms/step -
accuracy: 0.7385 - auc: 0.8014 - loss: 0.5565 - val_accuracy: 0.7409 - val_auc:
0.7997 - val_loss: 0.5509
Epoch 41/100
129/129      0s 3ms/step -
accuracy: 0.7418 - auc: 0.8025 - loss: 0.5539 - val_accuracy: 0.7530 - val_auc:
0.8020 - val_loss: 0.5481
Epoch 42/100
129/129      0s 2ms/step -
accuracy: 0.7434 - auc: 0.7922 - loss: 0.5653 - val_accuracy: 0.7538 - val_auc:
0.8000 - val_loss: 0.5477
Epoch 43/100
129/129      0s 3ms/step -
accuracy: 0.7414 - auc: 0.8029 - loss: 0.5522 - val_accuracy: 0.7441 - val_auc:
0.8020 - val_loss: 0.5463
Epoch 44/100
129/129      0s 2ms/step -
accuracy: 0.7374 - auc: 0.7961 - loss: 0.5564 - val_accuracy: 0.7498 - val_auc:
```

```
0.8009 - val_loss: 0.5428
Epoch 45/100
129/129          0s 2ms/step -
accuracy: 0.7440 - auc: 0.8020 - loss: 0.5511 - val_accuracy: 0.7441 - val_auc:
0.8028 - val_loss: 0.5419
Epoch 46/100
129/129          0s 2ms/step -
accuracy: 0.7498 - auc: 0.7989 - loss: 0.5553 - val_accuracy: 0.7457 - val_auc:
0.8015 - val_loss: 0.5435
Epoch 47/100
129/129          0s 2ms/step -
accuracy: 0.7426 - auc: 0.8029 - loss: 0.5536 - val_accuracy: 0.7474 - val_auc:
0.8007 - val_loss: 0.5423
Epoch 48/100
129/129          0s 2ms/step -
accuracy: 0.7476 - auc: 0.8057 - loss: 0.5465 - val_accuracy: 0.7441 - val_auc:
0.8024 - val_loss: 0.5424
Epoch 49/100
129/129          0s 2ms/step -
accuracy: 0.7327 - auc: 0.7870 - loss: 0.5675 - val_accuracy: 0.7457 - val_auc:
0.8032 - val_loss: 0.5410
Epoch 50/100
129/129          0s 2ms/step -
accuracy: 0.7504 - auc: 0.8022 - loss: 0.5516 - val_accuracy: 0.7433 - val_auc:
0.8040 - val_loss: 0.5396
Epoch 51/100
129/129          0s 3ms/step -
accuracy: 0.7370 - auc: 0.8030 - loss: 0.5501 - val_accuracy: 0.7457 - val_auc:
0.8032 - val_loss: 0.5392
Epoch 52/100
129/129          0s 2ms/step -
accuracy: 0.7237 - auc: 0.7910 - loss: 0.5603 - val_accuracy: 0.7506 - val_auc:
0.8007 - val_loss: 0.5411
Epoch 53/100
129/129          0s 2ms/step -
accuracy: 0.7385 - auc: 0.8000 - loss: 0.5501 - val_accuracy: 0.7433 - val_auc:
0.8018 - val_loss: 0.5398
Epoch 54/100
129/129          0s 2ms/step -
accuracy: 0.7388 - auc: 0.8002 - loss: 0.5503 - val_accuracy: 0.7474 - val_auc:
0.8036 - val_loss: 0.5385
Epoch 55/100
129/129          0s 3ms/step -
accuracy: 0.7334 - auc: 0.7950 - loss: 0.5543 - val_accuracy: 0.7433 - val_auc:
0.8059 - val_loss: 0.5377
Epoch 56/100
129/129          0s 2ms/step -
accuracy: 0.7366 - auc: 0.7951 - loss: 0.5537 - val_accuracy: 0.7514 - val_auc:
```

```
0.8037 - val_loss: 0.5365
Epoch 57/100
129/129          0s 3ms/step -
accuracy: 0.7451 - auc: 0.8043 - loss: 0.5446 - val_accuracy: 0.7482 - val_auc:
0.8045 - val_loss: 0.5379
Epoch 58/100
129/129          0s 2ms/step -
accuracy: 0.7536 - auc: 0.8019 - loss: 0.5458 - val_accuracy: 0.7441 - val_auc:
0.8030 - val_loss: 0.5389
Epoch 59/100
129/129          0s 2ms/step -
accuracy: 0.7440 - auc: 0.8017 - loss: 0.5465 - val_accuracy: 0.7466 - val_auc:
0.8022 - val_loss: 0.5385
Epoch 60/100
129/129          0s 2ms/step -
accuracy: 0.7461 - auc: 0.8055 - loss: 0.5432 - val_accuracy: 0.7474 - val_auc:
0.8020 - val_loss: 0.5384
Epoch 61/100
129/129          0s 2ms/step -
accuracy: 0.7478 - auc: 0.7955 - loss: 0.5520 - val_accuracy: 0.7498 - val_auc:
0.8025 - val_loss: 0.5365
Epoch 62/100
129/129          0s 2ms/step -
accuracy: 0.7403 - auc: 0.8049 - loss: 0.5412 - val_accuracy: 0.7409 - val_auc:
0.8028 - val_loss: 0.5375
Epoch 63/100
129/129          0s 3ms/step -
accuracy: 0.7502 - auc: 0.8044 - loss: 0.5425 - val_accuracy: 0.7360 - val_auc:
0.8025 - val_loss: 0.5388
Epoch 64/100
129/129          0s 3ms/step -
accuracy: 0.7433 - auc: 0.8024 - loss: 0.5457 - val_accuracy: 0.7474 - val_auc:
0.8012 - val_loss: 0.5380
Epoch 65/100
129/129          0s 2ms/step -
accuracy: 0.7434 - auc: 0.7924 - loss: 0.5558 - val_accuracy: 0.7466 - val_auc:
0.8023 - val_loss: 0.5373
Epoch 66/100
129/129          0s 2ms/step -
accuracy: 0.7452 - auc: 0.7996 - loss: 0.5489 - val_accuracy: 0.7498 - val_auc:
0.8036 - val_loss: 0.5358
Epoch 67/100
129/129          0s 2ms/step -
accuracy: 0.7414 - auc: 0.8007 - loss: 0.5445 - val_accuracy: 0.7466 - val_auc:
0.8043 - val_loss: 0.5338
Epoch 68/100
129/129          0s 3ms/step -
accuracy: 0.7447 - auc: 0.8056 - loss: 0.5431 - val_accuracy: 0.7498 - val_auc:
```

```
0.8046 - val_loss: 0.5339
Epoch 69/100
129/129          0s 3ms/step -
accuracy: 0.7257 - auc: 0.7897 - loss: 0.5606 - val_accuracy: 0.7417 - val_auc:
0.8062 - val_loss: 0.5336
Epoch 70/100
129/129          0s 3ms/step -
accuracy: 0.7463 - auc: 0.8075 - loss: 0.5392 - val_accuracy: 0.7393 - val_auc:
0.8046 - val_loss: 0.5358
Epoch 71/100
129/129          0s 3ms/step -
accuracy: 0.7475 - auc: 0.8099 - loss: 0.5392 - val_accuracy: 0.7417 - val_auc:
0.8033 - val_loss: 0.5359
Epoch 72/100
129/129          0s 3ms/step -
accuracy: 0.7410 - auc: 0.8060 - loss: 0.5431 - val_accuracy: 0.7530 - val_auc:
0.8039 - val_loss: 0.5353
Epoch 73/100
129/129          1s 5ms/step -
accuracy: 0.7375 - auc: 0.7962 - loss: 0.5507 - val_accuracy: 0.7506 - val_auc:
0.8033 - val_loss: 0.5345
Epoch 74/100
129/129          1s 5ms/step -
accuracy: 0.7512 - auc: 0.8006 - loss: 0.5453 - val_accuracy: 0.7449 - val_auc:
0.8034 - val_loss: 0.5343
Epoch 75/100
129/129          0s 3ms/step -
accuracy: 0.7492 - auc: 0.8092 - loss: 0.5373 - val_accuracy: 0.7425 - val_auc:
0.8056 - val_loss: 0.5329
Epoch 76/100
129/129          0s 3ms/step -
accuracy: 0.7409 - auc: 0.8032 - loss: 0.5440 - val_accuracy: 0.7490 - val_auc:
0.8041 - val_loss: 0.5342
Epoch 77/100
129/129          0s 3ms/step -
accuracy: 0.7492 - auc: 0.8030 - loss: 0.5417 - val_accuracy: 0.7498 - val_auc:
0.8035 - val_loss: 0.5351
Epoch 78/100
129/129          0s 3ms/step -
accuracy: 0.7326 - auc: 0.7904 - loss: 0.5542 - val_accuracy: 0.7498 - val_auc:
0.8046 - val_loss: 0.5339
Epoch 79/100
129/129          0s 3ms/step -
accuracy: 0.7454 - auc: 0.7999 - loss: 0.5492 - val_accuracy: 0.7466 - val_auc:
0.8048 - val_loss: 0.5346
Epoch 80/100
129/129          0s 3ms/step -
accuracy: 0.7376 - auc: 0.7989 - loss: 0.5476 - val_accuracy: 0.7490 - val_auc:
```

```
0.8074 - val_loss: 0.5332
Epoch 81/100
129/129          0s 3ms/step -
accuracy: 0.7421 - auc: 0.7968 - loss: 0.5494 - val_accuracy: 0.7449 - val_auc:
0.8021 - val_loss: 0.5359
Epoch 82/100
129/129          0s 3ms/step -
accuracy: 0.7403 - auc: 0.8000 - loss: 0.5446 - val_accuracy: 0.7530 - val_auc:
0.8030 - val_loss: 0.5329
Epoch 83/100
129/129          0s 3ms/step -
accuracy: 0.7373 - auc: 0.7967 - loss: 0.5507 - val_accuracy: 0.7490 - val_auc:
0.8045 - val_loss: 0.5326
Epoch 84/100
129/129          0s 3ms/step -
accuracy: 0.7517 - auc: 0.8114 - loss: 0.5310 - val_accuracy: 0.7498 - val_auc:
0.8047 - val_loss: 0.5335
Epoch 85/100
129/129          0s 2ms/step -
accuracy: 0.7339 - auc: 0.7912 - loss: 0.5524 - val_accuracy: 0.7506 - val_auc:
0.8033 - val_loss: 0.5341
Epoch 86/100
129/129          0s 3ms/step -
accuracy: 0.7443 - auc: 0.8045 - loss: 0.5427 - val_accuracy: 0.7417 - val_auc:
0.8038 - val_loss: 0.5347
Epoch 87/100
129/129          0s 3ms/step -
accuracy: 0.7386 - auc: 0.7966 - loss: 0.5497 - val_accuracy: 0.7490 - val_auc:
0.8044 - val_loss: 0.5319
Epoch 88/100
129/129          0s 3ms/step -
accuracy: 0.7415 - auc: 0.8087 - loss: 0.5362 - val_accuracy: 0.7514 - val_auc:
0.8045 - val_loss: 0.5327
Epoch 89/100
129/129          0s 3ms/step -
accuracy: 0.7510 - auc: 0.8095 - loss: 0.5336 - val_accuracy: 0.7457 - val_auc:
0.8058 - val_loss: 0.5311
Epoch 90/100
129/129          0s 3ms/step -
accuracy: 0.7391 - auc: 0.8014 - loss: 0.5439 - val_accuracy: 0.7482 - val_auc:
0.8071 - val_loss: 0.5303
Epoch 91/100
129/129          0s 3ms/step -
accuracy: 0.7513 - auc: 0.8122 - loss: 0.5316 - val_accuracy: 0.7474 - val_auc:
0.8060 - val_loss: 0.5323
Epoch 92/100
129/129          0s 2ms/step -
accuracy: 0.7331 - auc: 0.8017 - loss: 0.5413 - val_accuracy: 0.7506 - val_auc:
```

```

0.8059 - val_loss: 0.5318
Epoch 93/100
129/129          0s 2ms/step -
accuracy: 0.7545 - auc: 0.8184 - loss: 0.5257 - val_accuracy: 0.7482 - val_auc:
0.8064 - val_loss: 0.5307
Epoch 94/100
129/129          0s 2ms/step -
accuracy: 0.7341 - auc: 0.8017 - loss: 0.5451 - val_accuracy: 0.7506 - val_auc:
0.8057 - val_loss: 0.5299
Epoch 95/100
129/129          0s 3ms/step -
accuracy: 0.7411 - auc: 0.8026 - loss: 0.5441 - val_accuracy: 0.7482 - val_auc:
0.8044 - val_loss: 0.5335
Epoch 96/100
129/129          0s 3ms/step -
accuracy: 0.7506 - auc: 0.8022 - loss: 0.5425 - val_accuracy: 0.7514 - val_auc:
0.8066 - val_loss: 0.5303
Epoch 97/100
129/129          0s 3ms/step -
accuracy: 0.7429 - auc: 0.8043 - loss: 0.5425 - val_accuracy: 0.7522 - val_auc:
0.8049 - val_loss: 0.5305
Epoch 98/100
129/129          0s 3ms/step -
accuracy: 0.7365 - auc: 0.8003 - loss: 0.5461 - val_accuracy: 0.7474 - val_auc:
0.8096 - val_loss: 0.5286
Epoch 99/100
129/129          0s 2ms/step -
accuracy: 0.7412 - auc: 0.7982 - loss: 0.5472 - val_accuracy: 0.7490 - val_auc:
0.8076 - val_loss: 0.5293
Epoch 100/100
129/129          0s 3ms/step -
accuracy: 0.7347 - auc: 0.7912 - loss: 0.5545 - val_accuracy: 0.7466 - val_auc:
0.8070 - val_loss: 0.5296

--- Final Evaluation -High-Dropout Stability ---
Train Loss: 0.5194
Train Accuracy: 0.7535
Train AUC: 0.8204

Test Loss: 0.5286
Test Accuracy: 0.7474
Test AUC: 0.8096

```

[65]: # 3. Deep Dropout-Enhanced Neural Network

```
# 1. Early stopping to prevent overfitting
early_stop = callbacks.EarlyStopping(
```

```

        monitor='val_loss',
        patience=10,
        restore_best_weights=True
    )

# 2. Define the model
model = keras.Sequential([
    layers.Input(shape=(X_res.shape[1],)),
    layers.Dense(128, activation='relu'),
    layers.BatchNormalization(),
    layers.Dropout(0.4),

    layers.Dense(64, activation='relu'),
    layers.BatchNormalization(),
    layers.Dropout(0.3),

    layers.Dense(32, activation='relu'),
    layers.BatchNormalization(),
    layers.Dropout(0.2),

    layers.Dense(1, activation='sigmoid')
])

# 3. Compile the model
model.compile(
    optimizer=keras.optimizers.Adam(learning_rate=0.001),
    loss='binary_crossentropy',
    metrics=['accuracy', keras.metrics.AUC(name='auc')]
)

# 4. Train the model
history = model.fit(
    X_res, y_res,
    validation_data=(X_test, y_test),
    epochs=100,
    batch_size=32,
    callbacks=[early_stop],
    verbose=1
)

# 5. Evaluate the model
train_loss, train_accuracy, train_auc = model.evaluate(X_res, y_res, verbose=0)
test_loss, test_accuracy, test_auc = model.evaluate(X_test, y_test, verbose=0)

# 6. Print basic results
print("\n--- Final Evaluation -Deep Dropout-Enhanced Neural Network---")
print(f"Train Loss: {train_loss:.4f}")

```

```

print(f"Train Accuracy: {train_accuracy:.4f}")
print(f"Train AUC: {train_auc:.4f}")

print(f"\nTest Loss: {test_loss:.4f}")
print(f"Test Accuracy: {test_accuracy:.4f}")
print(f"Test AUC: {test_auc:.4f}")

# Save predictions for visualization
y_pred_prob = model.predict(X_test)
y_pred = (y_pred_prob > 0.5).astype(int)

```

Epoch 1/100
129/129 4s 6ms/step -
accuracy: 0.5929 - auc: 0.6308 - loss: 0.7351 - val_accuracy: 0.7231 - val_auc:
0.7653 - val_loss: 0.6185
Epoch 2/100
129/129 1s 4ms/step -
accuracy: 0.6587 - auc: 0.7151 - loss: 0.6440 - val_accuracy: 0.7304 - val_auc:
0.7780 - val_loss: 0.5765
Epoch 3/100
129/129 1s 4ms/step -
accuracy: 0.7022 - auc: 0.7502 - loss: 0.6033 - val_accuracy: 0.7417 - val_auc:
0.7920 - val_loss: 0.5527
Epoch 4/100
129/129 1s 4ms/step -
accuracy: 0.6970 - auc: 0.7532 - loss: 0.5961 - val_accuracy: 0.7482 - val_auc:
0.7980 - val_loss: 0.5390
Epoch 5/100
129/129 1s 4ms/step -
accuracy: 0.7047 - auc: 0.7614 - loss: 0.5896 - val_accuracy: 0.7425 - val_auc:
0.7997 - val_loss: 0.5345
Epoch 6/100
129/129 1s 4ms/step -
accuracy: 0.7191 - auc: 0.7840 - loss: 0.5597 - val_accuracy: 0.7498 - val_auc:
0.8036 - val_loss: 0.5296
Epoch 7/100
129/129 0s 3ms/step -
accuracy: 0.7242 - auc: 0.7861 - loss: 0.5573 - val_accuracy: 0.7506 - val_auc:
0.8015 - val_loss: 0.5302
Epoch 8/100
129/129 0s 3ms/step -
accuracy: 0.7200 - auc: 0.7808 - loss: 0.5638 - val_accuracy: 0.7538 - val_auc:
0.8051 - val_loss: 0.5281
Epoch 9/100
129/129 0s 3ms/step -
accuracy: 0.7386 - auc: 0.8003 - loss: 0.5400 - val_accuracy: 0.7530 - val_auc:
0.8039 - val_loss: 0.5270
Epoch 10/100

```
129/129          0s 3ms/step -
accuracy: 0.7519 - auc: 0.8059 - loss: 0.5334 - val_accuracy: 0.7449 - val_auc:
0.8038 - val_loss: 0.5299
Epoch 11/100
129/129          0s 3ms/step -
accuracy: 0.7338 - auc: 0.7958 - loss: 0.5438 - val_accuracy: 0.7530 - val_auc:
0.8061 - val_loss: 0.5272
Epoch 12/100
129/129          0s 3ms/step -
accuracy: 0.7322 - auc: 0.7901 - loss: 0.5496 - val_accuracy: 0.7538 - val_auc:
0.8054 - val_loss: 0.5256
Epoch 13/100
129/129          1s 4ms/step -
accuracy: 0.7362 - auc: 0.7894 - loss: 0.5514 - val_accuracy: 0.7530 - val_auc:
0.8081 - val_loss: 0.5238
Epoch 14/100
129/129          0s 3ms/step -
accuracy: 0.7329 - auc: 0.7912 - loss: 0.5467 - val_accuracy: 0.7555 - val_auc:
0.8068 - val_loss: 0.5247
Epoch 15/100
129/129          0s 3ms/step -
accuracy: 0.7358 - auc: 0.7999 - loss: 0.5387 - val_accuracy: 0.7506 - val_auc:
0.8107 - val_loss: 0.5213
Epoch 16/100
129/129          0s 3ms/step -
accuracy: 0.7325 - auc: 0.7903 - loss: 0.5456 - val_accuracy: 0.7522 - val_auc:
0.8069 - val_loss: 0.5238
Epoch 17/100
129/129          0s 4ms/step -
accuracy: 0.7296 - auc: 0.7891 - loss: 0.5489 - val_accuracy: 0.7522 - val_auc:
0.8067 - val_loss: 0.5214
Epoch 18/100
129/129          1s 4ms/step -
accuracy: 0.7368 - auc: 0.7947 - loss: 0.5438 - val_accuracy: 0.7490 - val_auc:
0.8076 - val_loss: 0.5226
Epoch 19/100
129/129          1s 5ms/step -
accuracy: 0.7507 - auc: 0.8141 - loss: 0.5246 - val_accuracy: 0.7482 - val_auc:
0.8081 - val_loss: 0.5233
Epoch 20/100
129/129          1s 4ms/step -
accuracy: 0.7438 - auc: 0.7948 - loss: 0.5420 - val_accuracy: 0.7498 - val_auc:
0.8088 - val_loss: 0.5218
Epoch 21/100
129/129          1s 4ms/step -
accuracy: 0.7347 - auc: 0.7937 - loss: 0.5458 - val_accuracy: 0.7530 - val_auc:
0.8087 - val_loss: 0.5237
Epoch 22/100
```

```
129/129          1s 4ms/step -
accuracy: 0.7449 - auc: 0.8069 - loss: 0.5304 - val_accuracy: 0.7571 - val_auc:
0.8077 - val_loss: 0.5230
Epoch 23/100
129/129          0s 3ms/step -
accuracy: 0.7464 - auc: 0.8041 - loss: 0.5332 - val_accuracy: 0.7498 - val_auc:
0.8084 - val_loss: 0.5224
Epoch 24/100
129/129          0s 3ms/step -
accuracy: 0.7400 - auc: 0.7984 - loss: 0.5396 - val_accuracy: 0.7498 - val_auc:
0.8107 - val_loss: 0.5210
Epoch 25/100
129/129          0s 3ms/step -
accuracy: 0.7392 - auc: 0.8003 - loss: 0.5386 - val_accuracy: 0.7555 - val_auc:
0.8084 - val_loss: 0.5203
Epoch 26/100
129/129          0s 3ms/step -
accuracy: 0.7475 - auc: 0.8087 - loss: 0.5272 - val_accuracy: 0.7538 - val_auc:
0.8110 - val_loss: 0.5191
Epoch 27/100
129/129          0s 3ms/step -
accuracy: 0.7396 - auc: 0.8096 - loss: 0.5328 - val_accuracy: 0.7522 - val_auc:
0.8137 - val_loss: 0.5168
Epoch 28/100
129/129          0s 3ms/step -
accuracy: 0.7508 - auc: 0.8081 - loss: 0.5275 - val_accuracy: 0.7579 - val_auc:
0.8108 - val_loss: 0.5190
Epoch 29/100
129/129          1s 4ms/step -
accuracy: 0.7318 - auc: 0.7890 - loss: 0.5498 - val_accuracy: 0.7522 - val_auc:
0.8122 - val_loss: 0.5162
Epoch 30/100
129/129          1s 4ms/step -
accuracy: 0.7461 - auc: 0.8199 - loss: 0.5164 - val_accuracy: 0.7579 - val_auc:
0.8131 - val_loss: 0.5173
Epoch 31/100
129/129          1s 4ms/step -
accuracy: 0.7370 - auc: 0.8009 - loss: 0.5346 - val_accuracy: 0.7563 - val_auc:
0.8127 - val_loss: 0.5175
Epoch 32/100
129/129          1s 4ms/step -
accuracy: 0.7451 - auc: 0.8082 - loss: 0.5283 - val_accuracy: 0.7595 - val_auc:
0.8132 - val_loss: 0.5155
Epoch 33/100
129/129          1s 4ms/step -
accuracy: 0.7443 - auc: 0.8042 - loss: 0.5327 - val_accuracy: 0.7530 - val_auc:
0.8114 - val_loss: 0.5173
Epoch 34/100
```

```
129/129          1s 4ms/step -
accuracy: 0.7461 - auc: 0.8006 - loss: 0.5360 - val_accuracy: 0.7538 - val_auc:
0.8111 - val_loss: 0.5167
Epoch 35/100
129/129          1s 5ms/step -
accuracy: 0.7477 - auc: 0.8037 - loss: 0.5316 - val_accuracy: 0.7530 - val_auc:
0.8126 - val_loss: 0.5159
Epoch 36/100
129/129          1s 4ms/step -
accuracy: 0.7536 - auc: 0.8125 - loss: 0.5251 - val_accuracy: 0.7595 - val_auc:
0.8152 - val_loss: 0.5141
Epoch 37/100
129/129          1s 4ms/step -
accuracy: 0.7478 - auc: 0.8028 - loss: 0.5330 - val_accuracy: 0.7530 - val_auc:
0.8140 - val_loss: 0.5142
Epoch 38/100
129/129          1s 4ms/step -
accuracy: 0.7498 - auc: 0.8097 - loss: 0.5263 - val_accuracy: 0.7522 - val_auc:
0.8170 - val_loss: 0.5135
Epoch 39/100
129/129          0s 3ms/step -
accuracy: 0.7454 - auc: 0.8076 - loss: 0.5293 - val_accuracy: 0.7595 - val_auc:
0.8173 - val_loss: 0.5127
Epoch 40/100
129/129          0s 3ms/step -
accuracy: 0.7500 - auc: 0.8055 - loss: 0.5294 - val_accuracy: 0.7547 - val_auc:
0.8140 - val_loss: 0.5154
Epoch 41/100
129/129          0s 4ms/step -
accuracy: 0.7417 - auc: 0.8038 - loss: 0.5329 - val_accuracy: 0.7579 - val_auc:
0.8187 - val_loss: 0.5097
Epoch 42/100
129/129          0s 3ms/step -
accuracy: 0.7369 - auc: 0.8090 - loss: 0.5280 - val_accuracy: 0.7571 - val_auc:
0.8182 - val_loss: 0.5111
Epoch 43/100
129/129          0s 3ms/step -
accuracy: 0.7445 - auc: 0.8074 - loss: 0.5289 - val_accuracy: 0.7514 - val_auc:
0.8188 - val_loss: 0.5121
Epoch 44/100
129/129          0s 3ms/step -
accuracy: 0.7472 - auc: 0.8113 - loss: 0.5232 - val_accuracy: 0.7587 - val_auc:
0.8179 - val_loss: 0.5105
Epoch 45/100
129/129          0s 3ms/step -
accuracy: 0.7540 - auc: 0.8180 - loss: 0.5180 - val_accuracy: 0.7522 - val_auc:
0.8177 - val_loss: 0.5113
Epoch 46/100
```

```
129/129          0s 4ms/step -
accuracy: 0.7470 - auc: 0.8132 - loss: 0.5211 - val_accuracy: 0.7530 - val_auc:
0.8166 - val_loss: 0.5119
Epoch 47/100
129/129          1s 4ms/step -
accuracy: 0.7318 - auc: 0.8079 - loss: 0.5280 - val_accuracy: 0.7611 - val_auc:
0.8195 - val_loss: 0.5080
Epoch 48/100
129/129          0s 4ms/step -
accuracy: 0.7545 - auc: 0.8194 - loss: 0.5131 - val_accuracy: 0.7587 - val_auc:
0.8220 - val_loss: 0.5062
Epoch 49/100
129/129          0s 4ms/step -
accuracy: 0.7602 - auc: 0.8234 - loss: 0.5085 - val_accuracy: 0.7587 - val_auc:
0.8222 - val_loss: 0.5047
Epoch 50/100
129/129          0s 4ms/step -
accuracy: 0.7556 - auc: 0.8180 - loss: 0.5138 - val_accuracy: 0.7595 - val_auc:
0.8207 - val_loss: 0.5071
Epoch 51/100
129/129          0s 4ms/step -
accuracy: 0.7350 - auc: 0.7990 - loss: 0.5351 - val_accuracy: 0.7555 - val_auc:
0.8214 - val_loss: 0.5050
Epoch 52/100
129/129          0s 4ms/step -
accuracy: 0.7487 - auc: 0.8114 - loss: 0.5220 - val_accuracy: 0.7595 - val_auc:
0.8219 - val_loss: 0.5059
Epoch 53/100
129/129          1s 4ms/step -
accuracy: 0.7417 - auc: 0.8051 - loss: 0.5313 - val_accuracy: 0.7547 - val_auc:
0.8203 - val_loss: 0.5073
Epoch 54/100
129/129          1s 4ms/step -
accuracy: 0.7438 - auc: 0.8199 - loss: 0.5147 - val_accuracy: 0.7571 - val_auc:
0.8220 - val_loss: 0.5048
Epoch 55/100
129/129          0s 4ms/step -
accuracy: 0.7438 - auc: 0.8028 - loss: 0.5342 - val_accuracy: 0.7538 - val_auc:
0.8224 - val_loss: 0.5054
Epoch 56/100
129/129          0s 3ms/step -
accuracy: 0.7478 - auc: 0.8134 - loss: 0.5193 - val_accuracy: 0.7514 - val_auc:
0.8224 - val_loss: 0.5083
Epoch 57/100
129/129          0s 4ms/step -
accuracy: 0.7512 - auc: 0.8062 - loss: 0.5282 - val_accuracy: 0.7563 - val_auc:
0.8238 - val_loss: 0.5053
Epoch 58/100
```

```
129/129      1s 4ms/step -
accuracy: 0.7442 - auc: 0.8115 - loss: 0.5229 - val_accuracy: 0.7514 - val_auc:
0.8251 - val_loss: 0.5029
Epoch 59/100
129/129      1s 4ms/step -
accuracy: 0.7483 - auc: 0.8150 - loss: 0.5206 - val_accuracy: 0.7587 - val_auc:
0.8262 - val_loss: 0.5009
Epoch 60/100
129/129      1s 4ms/step -
accuracy: 0.7621 - auc: 0.8239 - loss: 0.5086 - val_accuracy: 0.7603 - val_auc:
0.8228 - val_loss: 0.5090
Epoch 61/100
129/129      1s 4ms/step -
accuracy: 0.7407 - auc: 0.8100 - loss: 0.5271 - val_accuracy: 0.7555 - val_auc:
0.8275 - val_loss: 0.5016
Epoch 62/100
129/129      1s 4ms/step -
accuracy: 0.7533 - auc: 0.8146 - loss: 0.5212 - val_accuracy: 0.7579 - val_auc:
0.8278 - val_loss: 0.5012
Epoch 63/100
129/129      1s 4ms/step -
accuracy: 0.7590 - auc: 0.8141 - loss: 0.5192 - val_accuracy: 0.7579 - val_auc:
0.8253 - val_loss: 0.5032
Epoch 64/100
129/129      1s 4ms/step -
accuracy: 0.7597 - auc: 0.8214 - loss: 0.5132 - val_accuracy: 0.7538 - val_auc:
0.8246 - val_loss: 0.5051
Epoch 65/100
129/129      1s 4ms/step -
accuracy: 0.7537 - auc: 0.8204 - loss: 0.5137 - val_accuracy: 0.7595 - val_auc:
0.8252 - val_loss: 0.5013
Epoch 66/100
129/129      1s 4ms/step -
accuracy: 0.7582 - auc: 0.8231 - loss: 0.5111 - val_accuracy: 0.7538 - val_auc:
0.8246 - val_loss: 0.5017
Epoch 67/100
129/129      1s 4ms/step -
accuracy: 0.7501 - auc: 0.8095 - loss: 0.5238 - val_accuracy: 0.7563 - val_auc:
0.8262 - val_loss: 0.5013
Epoch 68/100
129/129      1s 4ms/step -
accuracy: 0.7458 - auc: 0.8045 - loss: 0.5270 - val_accuracy: 0.7538 - val_auc:
0.8236 - val_loss: 0.5037
Epoch 69/100
129/129      1s 4ms/step -
accuracy: 0.7298 - auc: 0.7998 - loss: 0.5402 - val_accuracy: 0.7522 - val_auc:
0.8264 - val_loss: 0.5023
```

```

--- Final Evaluation -Deep Dropout-Enhanced Neural Network---
Train Loss: 0.4959
Train Accuracy: 0.7630
Train AUC: 0.8321

Test Loss: 0.5009
Test Accuracy: 0.7587
Test AUC: 0.8262
39/39          0s 4ms/step

```

[66]: # --- PART 2: Visualization shown for Best Neural Network---

```

# 1. Training vs Validation Loss and Accuracy
fig, axes = plt.subplots(1, 2, figsize=(14,5))

# Plot Loss
axes[0].plot(history.history['loss'], label='Training Loss')
axes[0].plot(history.history['val_loss'], label='Validation Loss')
axes[0].set_title('Loss over Epochs')
axes[0].set_xlabel('Epoch')
axes[0].set_ylabel('Loss')
axes[0].legend()

# Plot Accuracy
axes[1].plot(history.history['accuracy'], label='Training Accuracy')
axes[1].plot(history.history['val_accuracy'], label='Validation Accuracy')
axes[1].set_title('Accuracy over Epochs')
axes[1].set_xlabel('Epoch')
axes[1].set_ylabel('Accuracy')
axes[1].legend()

plt.tight_layout()
plt.show()

# Calculate confusion matrix
cm = confusion_matrix(y_test, y_pred)
# Calculate ROC
fpr, tpr, thresholds = roc_curve(y_test, y_pred_prob)
roc_auc = auc(fpr, tpr)
# Create side-by-side plots
fig, axes = plt.subplots(1, 2, figsize=(14,6))

# Confusion Matrix Plot
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', ax=axes[0])
axes[0].set_title('Confusion Matrix')
axes[0].set_xlabel('Predicted Labels')
axes[0].set_ylabel('True Labels')

```

```

# ROC Curve Plot
axes[1].plot(fpr, tpr, label=f'ROC Curve (AUC = {roc_auc:.2f})')
axes[1].plot([0, 1], [0, 1], linestyle='--', color='gray')
axes[1].set_xlabel('False Positive Rate')
axes[1].set_ylabel('True Positive Rate')
axes[1].set_title('Receiver Operating Characteristic (ROC) Curve')
axes[1].legend(loc='lower right')

plt.tight_layout()
plt.show()

# 4. Learning Curve with Generalization Gap
fig, ax = plt.subplots(figsize=(8,5))
ax.plot(history.history['loss'], label='Training Loss')
ax.plot(history.history['val_loss'], label='Validation Loss')
ax.fill_between(
    range(len(history.history['loss'])),
    np.array(history.history['loss']),
    np.array(history.history['val_loss']),
    color='gray', alpha=0.3, label='Generalization Gap'
)
ax.set_title('Learning Curve with Generalization Gap')
ax.set_xlabel('Epoch')
ax.set_ylabel('Loss')
ax.legend()
plt.show()

# 5. Decision Boundary (only for 2D features)
if X_res.shape[1] == 2:
    from matplotlib.colors import ListedColormap

    x_min, x_max = X_res[:, 0].min() - 1, X_res[:, 0].max() + 1
    y_min, y_max = X_res[:, 1].min() - 1, X_res[:, 1].max() + 1
    xx, yy = np.meshgrid(np.arange(x_min, x_max, 0.01),
                         np.arange(y_min, y_max, 0.01))
    Z = model.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = (Z > 0.5).astype(int)
    Z = Z.reshape(xx.shape)



    plt.figure(figsize=(8,6))
    plt.contourf(xx, yy, Z, cmap=ListedColormap(('lightblue', 'lightcoral')))
    plt.scatter(X_res[:, 0], X_res[:, 1], c=y_res, cmap=ListedColormap(('blue', ↴
    'red')))
    plt.title('Decision Boundary')
    plt.xlabel('Feature 1')
    plt.ylabel('Feature 2')

```


```

    plt.show()
else:
    print("Decision boundary plot skipped (input features > 2D).")

```



Learning Curve with Generalization Gap



Decision boundary plot skipped (input features > 2D).